
ENTROPIC ATTRACTOR NETWORKS: A SCALE-AGNOSTIC, FIRST-PRINCIPLES ROUTE TO INFERENCE AND LEARNING IN BRAINS AND MACHINES

WHITE PAPER, VERSION 0.9 (DRAFT, NOT YET REVIEWED BY COLLEAGUES)

Tamas Spisak¹

Predictive Neuroscience Lab, Center for Translational Neuro- and Behavioral Sciences
University Medicine Essen, University of Duisburg-Essen

Saturday 3rd October, 2026

Contributors: Balint Kincses, Raviteja Kotikalapudi, Robert Englert, Jenya Lerer, Kevin Hoffschlag, Jialin Li, Giuseppe Gallitto (Predictive Neuroscience Lab).

Commenters and advisors (listing does not imply endorsement): none yet. Colleagues whose comments or advice shape future versions will be listed here, with their permission (see About this document).

Abstract


Artificial intelligence still lacks what biological brains do routinely: learning continually without overwriting old knowledge, representing uncertainty, and acting on an internal model on a small energy budget. Brains show that these problems can be solved, and neuroscience has revealed much about how. What is missing is a conceptual bridge at the fundamental level: we do not know how and why living matter, and brains in particular, came to have these capacities. This holds back AI, which lacks principles from which to build them, and neuroscience itself, whose detailed descriptions at the micro- and mesoscales are not yet tied together in a scale-agnostic account of where these capacities come from. Building the bridge requires a model class that is derived from physics rather than assumed, holds across spatial and temporal scales, and is at once a credible generative model of brain dynamics and a capable, trainable learner. This white paper sets out a research programme to build that bridge (see the graphical abstract below). From a few physical assumptions about coarse observables over hidden microstates, a minimal unit is constructed in which inference and learning arise as fast and slow entropic relaxation, not from designed rules. Composed into networks, such units form *entropic attractor networks* (eANNs), whose attractors are learned, approximately orthogonal priors about the external world (Line A). Fitted to fMRI, eANNs become generative models of large-scale brain dynamics with testable, mechanistic predictions (Line B); built in software and on emerging hardware, they offer an alternative approach to AI (Line C).

¹Correspondence to: tamas.spisak@uk-essen.de


Line A "Approach and Theory": first principles & falsifiable predictions

A physical account for perception and learning

Entropic (self-orthogonalizing) attractor networks



Spisak et al., in prep. 2026



Spisak & Friston, 2026

Learning

$$\Delta J_{ij} \propto \sigma_i \sigma_j - \underbrace{L_i(h_i + \sum_{k \neq i} J_{ik} \sigma_k)}_{\text{observed correlation (Hollman)}} - \underbrace{L_j(h_j + \sum_{k \neq j} J_{jk} \sigma_k)}_{\text{predicted correlation (anti-Bickham)}}$$

- predictive coding
- self-orthogonalization
- separates latent causes


Inference

$$\mathbb{E}_q[n_i] = L(h_i) = L \left(h_i + \sum_{j \neq i} J_{ij} \sigma_j \right)$$


- Bayesian belief update
- attractor states: priors
- MCMC sampling from the posterior
- continual and sequence learning

Line B "Brains": Computational generative brain modelling

Task and resting fMRI dynamics




Tasks-related activations as Bayesian inference



Englert et al., 2026; python package attractome

Line C "Computers": Emergent learning as an alternative route to AI

Bayesian inference via attractors



Spisak and Friston, 2026; Spisak et al. in prep. 2026

Contributors and collaborators: Balint Kincses, Raviteja Kotikalpudi, Robert Englert, Jenya Lerer, Jialin Li, Kevin Hoffschlag, Giuseppe Gallitto (Predictive Neuroscience Lab)

Commenters (listing requires permission and does not imply endorsement): N/A (see ‘About this document’)

About this document

This white paper sets down and shares a research vision before it has been fully formalized and tested.

It grew out of several years of empirical, simulation and theoretical work, much of it still unpublished. Taken together, these results suggest a general and, we think, intuitive picture of how inference and learning can emerge from suitably organized matter and how understanding these processes can contribute both to neuroscience and AI research. The purpose of this document is to state that picture explicitly, so that it can be developed rigorously and tested piece by piece.

Two peer-reviewed papers form its published core:

- [Spisak and Friston \[2026\]](#). Derives entropic, and specifically *self-orthogonalizing attractor networks* from first principles, without hand-designed learning or inference rules, and shows analytically and in simulations that they perform Bayesian inference with attractors as priors, develop approximately orthogonal attractors, learn sequences and resist catastrophic forgetting.
- [Englert et al. \[2026\]](#). Shows that such entropic attractor networks - built directly from human functional connectivity - indeed exhibit approximately orthogonal attractor-representations and provides initial evidence that the attractor landscape corresponds to canonical large-scale brain networks, replicates across datasets, and captures task-related changes and altered dynamics in pathological conditions.

We expect some of the ideas that we describe below in the proposed research programme to be wrong. We hope that most will survive, and we have tried to state each one so that its failure would itself be informative. The picture may turn out to unify several phenomena that are usually explained separately; that is a hope to be earned, not a claim.

What follows is a *research programme* in the sense of [Lakatos \[1978\]](#), rather than a finished theory: a small core of assumptions (Section 3.1), auxiliary hypotheses that can be revised as evidence comes in, and predictions that can fail (Section 7). Its direct value lies in the predictions it makes and the work it organizes, not in the breadth of what it currently explains.

To keep the status of each statement visible, key claims carry a tag:

Tag	Meaning
PUBLISHED	Shown in peer-reviewed work, mainly Spisak and Friston [2026] and Englert et al. [2026]
IN PREPARATION	Derived or simulated in work being prepared for publication
HYPOTHESIS	A heuristic expectation, partly supported by simulations or data, that the programme will formalize and test
SPECULATIVE	An appealing but distant possibility, flagged so that it is not mistaken for a claim

The document is versioned and will be updated as results arrive. Comments, criticism and collaboration ideas are welcome. We particularly welcome criticism from colleagues who work with other approaches,

including biophysical whole-brain models, task-optimized deep networks and the free energy principle. Serious objections will be recorded and answered in later versions, and readers who contribute will be acknowledged by name, with their permission. Contributors are listed on the title page, as are commenters and advisors whose comments or advice shaped a version; being listed records that input and does not imply endorsement of any claim. Responsibility for the content rests with the author.

Executive summary

The challenge. Biological brains learn continually without overwriting what they know, represent uncertainty, and act and plan based on internal models of the world, all on about 20 watts. These remain open problems for artificial intelligence. Neuroscience, in turn, has been revealing a great deal about the corresponding neural mechanisms, but has no account that explains, from first principles and across scales, where these capacities come from. Both problems share a root: a missing conceptual bridge between the physics of living matter and the computations it performs. Building it calls for a model class that is derived rather than designed, holds across scales, and is at once a credible generative model of brain dynamics and a capable, trainable learner.

The hypotheses. This white paper sets out a research programme to bridge this gap, built on three hypotheses, each tested by one line of work:

1. **Matter can infer and learn (Line A – Approach and theory).** In their most elementary form, inference and learning are capacities of suitably organized matter, not algorithms imposed on it. *Line A* focuses on working out a formal theory of this construction, and extends it to hierarchy, action and planning in complex networks of such adaptive physical units. The result is a broader family of networks we call *entropic attractor networks* (eANNs): stochastic attractor networks whose inference and learning rules are not designed but follow from the entropy-driven relaxation of coarse observables over hidden microstates.
2. **Biology exploits this across scales (Line B – Brains).** Life and biological intelligence use this elementary form of inference in many guises, stacked across spatial and temporal scales, and show quantifiable signatures of it. Evolution, on this view, may not have had to invent these capacities so much as find and refine them, shaping substrates that exploit, ever better, a law that suitably organized matter already obeys. *Line B* establishes the conditions by which eANNs can be fitted to fMRI data, to serve as a generative computational model of large-scale brain dynamics, tests the theory’s predictions, offers a principled test of which computations are carried by large-scale attractors and which by upstream, mesoscale processes, and probes the approach on a neuroscientific testbed focusing on learning, extinction, renewal, expectations and their effects on perception (e.g. on pain).
3. **Machines can do the same, natively (Line C – Computers).** Mainstream artificial intelligence has diverged from this elementary form of perception and learning, but machines can be built that display it natively, with perception and learning emerging from information geometric structure and thermodynamic primitives rather than being imposed by algorithms. *Line C* develops eANNs as an alternative route to AI, tests them on a neuroscience-inspired benchmark battery, and explores native implementation on thermodynamic, memristive, nanomagnetic and other emerging hardware, without imposing inference and learning rules explicitly.

The core idea (Ansatz). We start from parsimonious physical assumptions: timescale-separated, slow coarse observables over a system of simple, physically plausible microstates, and a distinguished point of view over the bounded coordinates they define. Then, we construct a minimal unit whose fast dynamics implement perception as inference and whose slow dynamics implement learning. The resulting minimal unit minimizes not only thermodynamic free energy but – owing to the natural normalization of the distributions belonging to the coarse observables – also exactly minimizes Friston’s variational free energy, licensing a Bayesian reading of the inference process. In the resulting information geometry – which we call the *self-aligning bundle* – a prediction error is a normalization-preserving and therefore expectation-centred tangent vector, and a synapse is the thermodynamic Legendre conjugate of a correlation (Section 3.2). Composed into networks, such units give rise to a special class of stochastic recurrent neural networks – eANNs, whose attractors store approximately orthogonal Bayesian priors, whose local learning rule takes a predictive-coding form, and which learn sequences and keep learning without catastrophic forgetting. Inference in these networks is fast

entropic relaxation, and learning is slow constrained maximum entropy, and their defining property is that learning self-orthogonalizes the attractors (Section 4). Noise in these networks is not a nuisance but encodes precision, and learned asymmetric couplings turn the landscape into directed, metastable sequence dynamics. Because the inferential form and the attractor landscape survive coarse-graining, models of the same family apply from single units (e.g. neurons) to larger, coarser modules (brain regions), with the substrate setting the particular member of the family (its reference measure, the bounds of its coordinates and the type of interaction between them): this is what makes the approach widely applicable, turning functional MRI a legitimate neuroimaging test bed (Section 4.2). Attractors, in this picture, are how a network creates new, slower timescales of its own, the same kind of timescale separation from which perception and learning were derived at the first place, so that the elementary step can repeat at the next level (Section 5).

What success looks like. This research programme does not promise a full account of the brain’s complexity, but a firmer grasp of its roots: predictions derived from one construction, about signatures shared by the whole family of models it defines, that hold in a scale-agnostic way both in brains and machines; generative models of large-scale brain dynamics whose components have a clear computational meaning; and learning machines whose continual learning and uncertainty handling are built in rather than added on. Each line may produce valuable insights even if the others stall, and all predictions are stated in advance with explicit criteria for failure (Section 7).

1. Problem statement

Brains solve problems that AI still struggles with, but we lack a conceptual bridge, at the fundamental level, from the physics of matter to these capacities.

Neuroscience has learned a great deal about how brains master situations that current mainstream approaches to artificial intelligence struggle with, like learning continually without overwriting old, but potentially still relevant knowledge, representing uncertainty and precision of evidence, and planning action based on an internal causal model, all this on a small energy budget. However neuroscience and biology has yet to explain how and why living matter, and brains in particular, came to have these capacities in the first place. This holds back AI and neuroscience alike. Below we describe the gap from both sides (1.1, 1.2), set the scope of what we aim to explain (1.3), and ask why existing frameworks do not close it on their own (1.4).

1.1 The asymmetry between artificial and biological intelligence

Modern AI writes, reasons and generates images, but its convergence towards biological intelligence is strikingly asymmetric. Large models are trained offline, with global loss functions, on vast data, using methods shaped by the constraints of the von Neumann hardware architecture, the “hardware lottery” [Hooker, 2021]. They struggle with capacities that brains show routinely:

- **Continual learning.** In brains, learning and inference are inseparable and never stop. Artificial networks are trained sequentially and tend to overwrite old knowledge (catastrophic forgetting). Long-horizon tasks show the same limitation within a single task: in a recent comparison, language-model agents given a day of compute fell behind the gains that repeated sampling would predict, whereas expert humans kept improving over two weeks [Mang et al., 2026] (an informal comparison, cited here as motivation rather than evidence; similar observations are common). From this programme’s perspective, this is less a matter of protecting old weights than of missing timescales: inference runs, but there is no slower level of structure that keeps adapting to the task while it is being performed, and no level slower still that consolidates what was learned (Sections 3, 4.3 and 5).
- **Uncertainty.** Brains act on calibrated uncertainty; most artificial networks represent it poorly or only through add-on machinery.
- **Frugality and locality.** Brains learn with local plasticity driven by noisy signals, on a tiny energy budget, whereas backpropagation needs globally orchestrated error transport and runs on deterministic hardware that spends much of its energy on suppressing noise.
- **Agency and planning.** Brains plan and act on an internal, actionable world model. In AI, despite recent achievements, this is still largely engineered on top.

Brains show that all of this is physically achievable, and neuroscience has identified many of the mechanisms involved, from local plasticity [Bi and Poo, 1998] and replay [Wilson and McNaughton, 1994] to attractor

dynamics [Khona and Fiete, 2022]. What AI lacks is therefore less an existence proof than a principled theoretical foundation: an account of why these mechanisms work that is general enough to carry over to other substrates.

1.2 A missing bridge at the fundamental level

Neuroscience describes the brain in remarkable detail, from molecules and synapses to circuits and whole-brain networks, and has models at each of these scales. What is still missing is an agreed, testable account - at the level of physics - of how living matter can infer and learn at all, and why similar capacities recur at very different scales, from single cells to whole brains [Lyon, 2005, Levin, 2019]. Candidate frameworks exist, from the free energy principle to basal cognition and physical learning, but none yet leads from first principles to a model that can be fitted to, and tested against, brain data.

One symptom is that models at different scales are hard to connect. Bottom-up approaches build whole-brain models from local biophysics (neural masses, neural fields, mean-field models). They are mechanistic, but the space of plausible variants is large, and validating them against macro-scale data is difficult. Top-down approaches describe macro-scale phenomena (functional connectivity, resting-state networks) without a mechanism that generates them. The two research cultures rarely meet [Rosas et al., 2025]. A bridge at the fundamental level would give both a common reference: a principle that holds at every scale, against which bottom-up mechanisms and top-down descriptions can be tested.

1.3 Scope: the roots of adaptive behaviour, not its full complexity

We do not doubt that the brain’s complexity will remain a major challenge for a long time. Its many cell types, its molecular machinery, its developmental and evolutionary history and its individual variability will not be derived from a handful of assumptions. Our claim is narrower: the roots from which this complexity grows can be understood much better than they are now. Behind the plurality of implementations that biology shows at different levels, from molecular networks to neural circuits, and across very different substrates, there may be shared principles of organization that all of them approximate, each in its own way. One possibility, which we find plausible but state as a hypothesis, concerns how evolution relates to these principles. Evolution may not have had to invent inference and learning from scratch, as an assembly of ad hoc rules. It may instead have started from a property that suitably organized matter already has, if only rarely and imperfectly when assembled at random: slow structural coordinates that, once exposed to a structured environment, tend to align with it (Section 3.2). Evolution would then have found this law rather than devised it, and spent its effort on substrates that exploit it ever better [Fields and Levin, 2020, Levin, 2023]. Fields et al. [2021] and Fields and Levin [2022] have argued, from a different starting point, that inference-like behaviour is generic to physical systems that interact with an environment. If inference and learning are elementary capacities of suitably organized matter, and if living systems elaborate them by stacking new timescales on old ones while keeping each level approximately self-contained (Section 5), then understanding the elementary step, and the rules by which it composes, would already be a substantial step for both neuroscience and AI. For neuroscience, it would give a principled account of why some of the brain’s core mechanisms take the form they do; for AI, principles to build on rather than mechanisms to copy. Biological and engineered cognition could then be studied as branches of a common root.

What the scope means for the hypotheses. The same scope shapes the hypotheses and predictions that follow. Most of them are not claims about the brain’s full machinery but about the characteristic signatures of the model class that grows from the proposed root: priors stored as attractors, approximately orthogonal attractors (self-orthogonalization), noise that encodes precision, asymmetric couplings that carry activity along the landscape rather than down it (self-solenoidization), the timing of a predictive-coding learning signal, and an inferential form that survives coarse-graining (Sections 4 and 5). If evolution has indeed settled on this class of solutions among the many conceivable ways to infer, learn and compute, these signatures should be found broadly: at different scales, in different substrates and under very different implementations, each a different member of the same family, approximating the shared law to its own degree. That makes them useful twice. For neuroscience, they are markers to look for and to test (Section 7), and their absence in systems that meet the assumptions would count against the programme. For AI, they point to the properties worth building into machines that exploit the same phenomenon, rather than imitating its biological implementation.

1.4 Why existing frameworks do not close the gap on their own

- **Task-optimized deep networks (neuroconnectionism).** This programme has given neuroscience its most successful models of sensory representations, and a productive way of comparing brains and machines on the same tasks [Richards et al., 2019, Doerig et al., 2023]. Its great strength is task fidelity: the models actually perform the functions they are meant to explain. For the question posed here, two limits remain. Neurobiological validity is partial: learning is mostly left to backpropagation, units and dynamics are idealized, and the models are rarely generative models of whole-brain dynamics. More fundamentally, architecture, objective and learning rule are chosen by the modeller and judged by their fit to data. The approach can therefore show which functions a brain-like system can learn, but it does not set out to explain why suitably organized matter comes to infer and learn in the first place. It offers no first-principles route from the physics of the substrate to the computation, and this route, we argue, is what both neuroscience and AI currently lack.
- **Biophysical whole-brain models.** They sit at the other end of the spectrum. Mechanistic and biophysically interpretable at their chosen scale, they are the most direct link between large-scale brain dynamics and physiology [Deco et al., 2008, Sanz Leon et al., 2013, Breakspear, 2017]. Their strength is biological plausibility, but they have little task fidelity: they are not designed to learn or to perform the functions whose neural basis they describe, which also limits how directly they can inform AI. They are, moreover, generically bottom-up: the local mechanism is taken as given at one privileged scale, and large-scale behaviour is derived from it. The approach proposed here adds a scale-agnostic perspective instead, seeking the same family of inferential forms at every level of description (Section 4.2). It is meant to complement these models, not to replace them.
- **The Free Energy Principle and active inference.** They offer the broadest first-principles account of self-organization as inference, but their generality comes at a price: the conditions under which the inferential reading holds are debated. The principle itself - as all principles - is not falsifiable; it is meant to be a heuristic framework to build on. Here we aim at an explicit route from physics to a model that is falsifiable - and we argue that the free energy principle provides the right heuristic for that.
- **Physical learning, energy-based models and other local-rule systems.** Hopfield networks and Boltzmann machines [Hopfield, 1982, Ackley et al., 1985], their modern successors, dense associative memories with very large storage capacity and a close link to attention in transformers [Krotov and Hopfield, 2016, 2020, Ramsauer et al., 2020], equilibrium propagation and coupled learning in physical networks [Scellier and Bengio, 2017, Stern et al., 2021], and cellular automata, classical and neural [Wolfram, 1984, Mordvintsev et al., 2020], all show how far memory, computation and learning can go with local interactions. In each, however, the energy function, the update rule or the training objective is chosen by the designer, often for capacity or for supervised tasks; dense associative memories and neural cellular automata, for example, are usually trained end to end by gradient descent rather than by a local rule. Few of these models are tied to brain data.
- **Large-scale attractor and energy-landscape models of brain dynamics.** A small but growing family of models treats whole-brain activity as movement on an attractor or energy landscape, in a tradition that goes back to Amit [1989]. It includes multistable and ghost-attractor regimes of connectome-based mean-field models [Deco and Jirsa, 2012, Golos et al., 2015], coupled-oscillator (Kuramoto-type) models [Cabral et al., 2011], pairwise maximum-entropy (Ising-type) energy landscapes fitted to fMRI [Watanabe et al., 2013, Ezaki et al., 2017, Gu et al., 2018], Hopfield-type networks built from mouse cortical connectivity [Fasoli et al., 2026] or estimated from MEG [Gosti et al., 2024], and our own functional connectome-based attractor networks of the human brain [Englert et al., 2026], which already rest on the self-orthogonalizing construction of Spisak and Friston [2026]. Attractor-state perspectives have also been proposed for the default mode network and the dynamics of ongoing experience [Poerio and Karapanagiotidis, 2025]. These models capture real features of brain dynamics, but most are phenomenological: couplings are taken from structural or functional connectivity or fitted to data, and dynamics and storage rules are chosen ad hoc rather than derived. Why brain dynamics should take this form, and what computation the attractors carry, is largely left open.

Each of these supplies part of the bridge: models that perform tasks, mechanisms grounded in physiology, a first-principles language, or learning rules native to physics. Neuroconnectionism and biophysical modelling mark the two ends of one spectrum from task fidelity without a first-principles origin to biological plausibility without task fidelity. What is still missing is a construction that connects them: derived from first principles, valid across scales, able to learn and perform tasks, and testable in brains and machines alike.

2. Desiderata

Before proposing a solution, we fix the criteria it must meet; the rest of this paper should be judged against them.

#	Criterion	What it requires
D1	Physically grounded	Inference and learning follow from stated physical assumptions about coarse observables, not from algorithms imposed from outside.
D2	Scale-agnostic	Rules of the same family apply, approximately, at any level of description (unit, region, network), so that the model class transfers across scales and coarse measurements such as fMRI are legitimate targets (Section 4.2). Ideally, the model also explains how new levels of description arise from within the system (Section 5).
D3	Local	Learning uses only quantities available at the unit: no global loss and no error transport.
D4	Constructive and generative	The model says not only what such a system is but also hints how to construct it. As a simulation or in hardware, it can be built from its stated rules, ideally running natively on physical substrates with noise as a resource rather than a nuisance. As a brain model, it is not only fitted to data but is also capable to generates data from the fitted model that were not measured, including the effects of counterfactual interventions.
D5	Trainable	The same model class learns tasks: sequences, continual learning, inference under uncertainty.
D6	Probabilistic	Uncertainty is represented natively, at both fast (states) and slow (structure) timescales.
D7	Falsifiable	Assumptions are explicit, each can fail, and the model makes quantitative predictions with stated failure criteria.
D8	Interpretable	Model components map onto computational roles (priors, likelihoods, prediction errors) and onto measurable quantities.

3. The Ansatz: from coarse physical dynamics to a minimal inferring and learning unit

Instead of deriving inference for the most complex systems we know, we derive it for the simplest building block, and leave complexity to composition.

This follows from the scope set in Section 1.3: the target is the root from which complexity grows, not complexity itself. How the building blocks compose, and how each level creates the slower timescales on which the next can build, is taken up in Sections 4 and 5.

This section states the construction and what it yields. Derivations are in the manuscript “From physical dynamics to perception and learning” (Spisak et al., in preparation; available on request, preprint forthcoming) **(IN PREPARATION)**.

3.1 Assumptions

The construction uses a small set of explicit assumptions, chosen by Occam’s razor: whenever several structures could support a result, the least elaborate one is taken.

1. **A normalized reference ensemble.** Stationary densities of the system are positive and regular enough to admit local exponential coordinates; environmental forcing acts as an exponential tilt of a reference state.
2. **Coarse observables with a double separation of timescales.** Microscopic variables relax fastest, retained mesoscopic coordinates next, and slow structural coordinates (couplings) much more slowly. The environment varies on an intermediate timescale.
3. **Bounded coordinates and a point of view.** Retained coordinates are bounded, as in any finite physical system. One of them, the exposed coordinate, is followed from its own point of view; the others form its boundary. No Markov blanket is required: shielding from the environment is a special case, not a precondition.
4. **Minimal interaction and least-structured response.** Near a generic operating point, only the leading boundary–exposed interaction is kept, and the resting law of the exposed coordinate is the maximum-entropy law on its range.
5. **Causal renewal (for networks).** Coupled units renew together, each forming its prediction from the preceding state of its boundary.

Each assumption is an empirical question for a given substrate, and each can fail.

3.2 What follows

- **Bayesian structure from normalization.** Normalizing the tilted ensemble over fast and slow coordinates turns the reference into a prior, the tilt into a likelihood and the normalizer into model evidence. No beliefs are assigned to microstates.
- **Prediction error as a tangent.** On the non-parametric statistical manifold of normalized densities [Pistone and Sempi, 1995], every tangent is centred: “actual minus expected”. Along the fast direction, this centered tangent is exactly the unit’s prediction error. It follows from preserving normalization alone.
- **A sigmoid from boundedness.** The unit’s emergent mean response is the Langevin function (tanh for a binary coordinate), the Legendre map of the bounded maximum-entropy family. No activation function is chosen by hand.
- **Perception as fast relaxation.** Mesoscale entropic relaxation moves the fast state towards the posterior, which is free-energy descent at fixed structure.
- **Learning as slow relaxation.** Timescale separation makes the expected correlations of mesoscale coordinates the Legendre duals of the slow structural coordinates. Under a dynamical form of Jaynes’ constrained maximum entropy [Jaynes, 1957], the structure climbs the average evidence: a contrastive, Hebbian minus anti-Hebbian learning signal.
- **Uncertainty from fluctuations.** At both timescales, fluctuations sample the corresponding posterior; for the fast state, variance equals gain (a fluctuation–dissipation relation). For the slow structure, this is the idea of synaptic sampling [Kappel et al., 2015].

In the unit’s own terms, the two centered scores that drive perception and learning are:

The construction in two sentences

Perception. A prediction error, geometrically, is nothing more than a normalization-preserving, and therefore expectation-centred, tangent vector on a non-parametric statistical manifold.

$$\varepsilon_\mu = \underbrace{\mu}_{\text{actual state}} - \underbrace{\mathbb{E}[\mu \mid s, W]}_{\text{expected state (prediction)}} \quad (1)$$

Learning. A synapse is the thermodynamic Legendre conjugate of a correlation.

$$\psi(W) = \sup_q \left\{ \underbrace{W}_{\text{structural parameter}} \underbrace{\mathbb{E}_q[s \mu]}_{\text{expected correlation}} - \underbrace{D_{\text{KL}}(q \parallel q_0)}_{\text{relative entropy}} \right\} \quad (2)$$

$$\varepsilon_W = \underbrace{s \mu}_{\text{observed correlation}} - \underbrace{\mathbb{E}[s \mu \mid W]}_{\text{expected correlation}} \quad (3)$$

Here ψ is the log-partition function (a free energy) and q_0 the reference law of the unconstrained dynamics. The structural parameter and the expected correlation form a Legendre pair, with relative entropy as the dual potential. Since $\partial\psi/\partial W = \mathbb{E}[s \mu \mid W]$, the learning error is the gradient of the dual objective $W \cdot s \mu - \psi(W)$: learning climbs it until expected and observed correlations agree, which is Jaynes’ constrained maximum entropy in dynamical form. The statistical manifold is the non-parametric exponential manifold of [Pistone and Sempi \[1995\]](#), whose tangent vectors are centred, “actual minus expected”. Synapse is meant in the general sense: any slow coupling between units.

The result is a sufficient, not necessary, route: a minimal and physically plausible formal model of how matter can be organized to infer and learn. It recovers, as derived consequences, the inference and learning parts of the FEP for this class of systems, without requiring the FEP’s general claims.

A family, not a single model. Strictly, the construction defines a family of units and networks, not one set of equations. Three ingredients come from the substrate and select a member of the family: the reference measure of the microstates, the bounds of the retained coordinates (which fix the response function, the Langevin function in general and tanh only for a binary coordinate), and the type of intrinsic interaction between mesoscopic coordinates. What all members share is the Bayesian structure, the centred prediction error, the Legendre-dual learning signal and the signatures that follow from them. When this white paper speaks of the same rules at different scales or in different substrates, it means rules of this family.

Imperfect starts, self-alignment and headroom. The construction does not require finely tuned structure. Structural coordinates may start far from any useful configuration; the slow dynamics then move them until expected and observed correlations agree, so that the structure aligns itself with the statistics it is exposed to. This is one reason for the name *self-aligning bundle*. How far the alignment can go is limited by the substrate: by the range, precision and timescales of the structural coordinates it offers. A physical system may thus obey the law only weakly, with little headroom for improvement, and evolution and development can be read as processes that build substrates with more headroom (Sections 1.3 and 11) [\(HYPOTHESIS\)](#).

Why “entropic”. The name refers to the dynamics, not only to their endpoint. Entropy is often read as describing where a system ends up, but a long tradition treats relative entropy as what drives the motion: fluxes driven by entropy gradients in irreversible thermodynamics [[Onsager, 1931](#)]; the Fokker–Planck equation as a gradient flow of free energy, that is, of relative entropy to the stationary law [[Jordan et al., 1998](#)]; and the split of dynamics into a reversible, energy-driven part and an irreversible, entropy-driven part [[Grmela and Öttinger, 1997](#)], which mirrors the solenoidal-plus-gradient decomposition used here (Section 4.5). In this sense, inference in the units is a relaxation that lowers relative entropy towards the posterior, learning is constrained maximum entropy, and the solenoidal part of the dynamics leaves relative entropy unchanged. As with entropic forces, what drives all three is the multiplicity of hidden microstates, not a designed interaction. The exact form of the relaxation, for example the metric in which it is a gradient flow, is part of the derivation in preparation. The programme of Entropic Dynamics, which derives dynamics, including quantum theory, from entropic inference [[Caticha, 2011](#)], shares this intuition but not the construction.

4. From units to networks: entropic attractor networks

Composed in a straightforward way, the minimal units form entropic attractor networks (eANNs): stochastic attractor networks whose attractors are Bayesian priors and whose learning rule self-orthogonalizes them [Spisak and Friston, 2026].

4.1 Inference and learning rules

When units are coupled, each unit’s boundary consists of the other units’ states. Each node then performs local Bayesian inference, combining its bias (prior) with weighted evidence from the other nodes through the Langevin function L , and adapts its couplings with a local rule that contrasts observed with predicted correlations:

$$\mathbb{E}_q[\sigma_i] = L\left(b_i + \sum_{j \neq i} J_{ij} \sigma_j\right), \quad \Delta J_{ij} \propto \sigma_i \sigma_j - L\left(b_i + \sum_{k \neq i} J_{ik} \sigma_k\right) \sigma_j \quad (4)$$

The first term of the learning rule is Hebbian (observed correlation), the second anti-Hebbian (predicted correlation): a predictive-coding rule. It belongs to the family of contrastive, local learning rules whose relation to inference has been worked out for predictive-coding networks [Millidge et al., 2022] and, in the symmetric equilibrium case, for Boltzmann machines [Ackley et al., 1985]. The network is a generalized, stochastic Hopfield network, a NESS-generalization of Boltzmann machines, whose noisy dynamics amount to Markov chain Monte Carlo sampling from the global posterior. The construction in Section 3 fixes the timing of the rule: the prediction is formed from the preceding network state and paired with the renewed one. In network terms, the two statements of Section 3.2 become concrete: each node’s prediction error is the centred tangent of its state, and each coupling J_{ij} is the thermodynamic Legendre conjugate of the correlation between nodes i and j , which it adjusts until observed and predicted correlations agree.

One-dimensional units, binary questions. The minimality assumption makes every node, and therefore every attractor, one-dimensional: its continuous-Bernoulli state expresses graded evidence for or against a single event. This is why, with zero biases, the sign-inverse of every attractor is also a fixed point. An attractor and its mirror image are one prior, read as evidence for or against the same latent cause.

4.2 Across scales: one model family and the same landscape at every resolution

Both the generality of the approach and its application to fMRI rest on what survives coarse-graining: the form of the construction, and the attractor landscape it produces Partly (IN PREPARATION) (HYPOTHESIS) for brain data.

One family of forms. Timescale separation is used twice: within the unit, where it separates perception from learning (Section 3.1), and between levels, where fast modes are enslaved by slow ones that carry the long-run dynamics [Haken, 1983, Carr, 1981, Kuramoto and Nakao, 2019, Jafarian et al., 2021]. In whole-brain models, this has been proposed, and demonstrated in simulations, as the mechanism by which low-dimensional structure emerges [Runfola et al., 2025]. If a coarse unit, such as a column or a brain region, exposes one macro-observable, a conjugate field acting on its transition ratios again yields an exponential family, with a response function and a susceptibility equal to its variance. The particular reference family need not survive averaging, but the architecture does: one exposed statistic, a conjugate field, fluctuation–dissipation, and a Bayesian reading when the macrostate represents an external cause. What changes with substrate and scale is which member of the family applies: the reference measure of the microstates, the bounds of the retained coordinates (and with them the shape of the response function, tanh only in the binary case), and the type of intrinsic interaction between mesoscopic coordinates. Single neurons, cortical columns and brain regions are therefore not expected to obey identical equations, but rules of one family, which share its signatures. Rules of this family can therefore apply from single neurons, where the learning rule resembles spike-timing-dependent plasticity [Bi and Poo, 1998] (HYPOTHESIS), to brain regions (D2). Making this rigorous for networks is a goal of Line A (A2).

The same landscape. An fMRI parcellation does not create a new network. It reports the activity of one neural system in fewer coordinates, the parcel means $y = C\sigma$: it is a camera, and different atlases are different cameras pointed at the same scene. The attractors $p\mu$ of the fine system are whole-brain patterns. Writing each as its regional profile $a\mu = Cp\mu$ plus a within-parcel residual $r\mu$, the one essential assumption

is homogeneity: on the attractors that matter, units within a parcel participate alike, which is exactly what functional parcellations are designed for [Gordon et al., 2014, Eickhoff et al., 2018]:

$$p_i^\mu = a_{\alpha(i)}^\mu + r_i^\mu, \quad \|\mathbf{r}^\mu\| \ll \|\mathbf{p}^\mu\| \quad (5)$$

Because the residuals average to zero within each parcel, fine inner products split exactly into a size-weighted coarse inner product and a residual term that homogeneity makes small:

$$\langle \mathbf{p}^\mu, \mathbf{p}^\nu \rangle = \sum_{\alpha} |P_{\alpha}| a_{\alpha}^{\mu} a_{\alpha}^{\nu} + \langle \mathbf{r}^\mu, \mathbf{r}^\nu \rangle \quad (6)$$

Three consequences follow. First, attractors that differ across parcels remain visible, those that differ only within parcels merge, and nothing new should appear; in nested atlases [Urchs et al., 2019], higher resolution splits coarse attractors into descendants that average back to them. Second, approximate orthogonality at the fine scale is inherited at the coarse scale, in the size-weighted parcel metric (analyses in unweighted or z-scored coordinates therefore use size-balanced atlases). Third, only a functional camera is useful: spatially scattered random parcels average structured patterns away, shrinking them roughly with the square root of parcel size until at most a global mode remains, whereas contiguous random parcels degrade gradually, the finest attractors first. The claim is not that a network can be fitted at any scale, but that meaningful coarse-grainings preserve the landscape and meaningless ones erase it [Meshulam et al., 2019, Friston et al., 2021, Machta et al., 2013, Transtrum et al., 2015].

Assumptions, limits and tests. Besides homogeneity, detecting the landscape with fitted networks assumes that the slow dynamics are organized by a modest number of attractors, and that the coarse network models the observed regional process; orthogonality is a prediction (P1, P2), not a prerequisite. Nothing assumes that a parcel is a large neuron, that couplings survive block averaging, or that coarse-graining is an exact renormalization map. After haemodynamic deconvolution, fMRI is a slow camera for slow patterns [Friston et al., 2014, Heitmann and Breakspear, 2018]. Averaging cannot recover distinctions it has discarded, and residual haemodynamic mixing can mimic mixtures of successive states, which simulations must rule out before such mixtures are read as prior mixing (P10). Nested atlases in population and precision fMRI [Gordon et al., 2017] make the argument directly testable: project fine attractors onto coarse atlases and match them; check that coarse attractors split into descendants; measure homogeneity as the between-parcel share of each attractor’s norm; compare inherited inner products with dimension-matched nulls; and run random-parcellation controls. The resolution range is fixed in advance by test–retest reliability, and matching is repeated within single participants (B1; P18–P20).

In one sentence

Fit models of the same family at every resolution; the attractors should be the same objects, seen through a coarser or finer camera.

4.3 Emergent properties

- **Attractors as priors** (PUBLISHED). Attractor states store the network’s priors; ongoing activity combines them with current input into a posterior. Local inference at the nodes gives rise to inference at the level of the whole network.
- **Self-orthogonalization, the key property** (PUBLISHED) theory, in-silico demonstrations, first evidence in human fMRI. Even for correlated inputs, the predictive-coding rule drives attractors towards approximate orthogonality: a new pattern is stored as its residual from what is already known. The network thereby approximates a Kanter–Sompolinsky projector network [Personnaz et al., 1986, Kanter and Sompolinsky, 1987], whose attractors coincide with the leading eigenvectors of the coupling matrix, but with a local learning rule. Offline routes towards similar couplings are known as well: Hebbian unlearning [Hopfield et al., 1983] and “dreaming”, which provably drives Hebbian couplings towards the projector [Fachechi et al., 2019] and can exploit correlations in the data [Serricchio et al., 2025]. In effect, self-orthogonalization is an on-line, approximate principal component analysis, in the tradition of Hebbian and anti-Hebbian subspace learning [Oja, 1982, Földiák, 1990, Pehlevan and Chklovskii, 2019]. Orthogonal attractors can act towards separating latent causes and use storage capacity efficiently. Furthermore, they serve as a basis for the emergent capabilities discussed below.

- **Attractor–prior duality** Symmetric case: formally established equivalence; antisymmetric extension **(HYPOTHESIS)**. This is the Hopfield network–restricted Boltzmann machine equivalence [Barra et al., 2012] in disguise. Approximate orthogonality makes the network statistically equivalent to a two-layer network in which the attractors are actual hidden nodes, connected only to the original nodes. Our heuristic is that, with an antisymmetric coupling component, non-equilibrium steady states and solenoidal dynamics appear as within-layer connections between the attractor nodes. The appeal is direct: attractors act as integrator nodes that pool evidence from many regions and feed the result back, performing (Bayesian) information integration. A change in a single connection can then be decomposed into changes in its contribution to individual attractors.
- **Superposition and prior mixing** **(HYPOTHESIS)** observed in simulations. Linear combinations of stored priors can exist as attractors or as “ghost attractors”, regions of slow dynamics without a true minimum [Deco and Jirsa, 2012]. Such slow regions can do real work: in classical and dense associative memories, patterns remain transiently retrievable above capacity from the slow remnants of their former basins [Clark, 2026]. Unlike the spurious states of classical Hopfield networks [Amit et al., 1985], they are useful here: the network can span the whole subspace of learned patterns and generalize to new inputs. Dense associative memories can likewise combine exponential capacity with emergent, useful minima that are not stored patterns [Hoover et al., 2025]. Mixing can also unfold in time, as quasi-periodic alternation across several attractors whose time-averaged posterior matches the data. Self-orthogonalization is what makes this work, and it offers a model for oscillatory brain dynamics [Aitchison and Lengyel, 2016, Echeveste et al., 2020].
- **Self-solenoidization** **(HYPOTHESIS)** observed in simulations. With our emergent learning rule, sequential presentation yields asymmetric couplings [Sompolinsky and Kanter, 1986, Kleinfeld, 1986]. During transitions, the landscape set by the symmetric couplings tends to absorb the flow components along the free-energy gradient, via a mechanism closely related to self-orthogonalization. What remains is encoded mainly in the antisymmetric couplings, as a transversal flow along the free-energy isocontours: approximately solenoidal flow that replays sequences without changing the steady state (transversality derivation is in progress, full solenoidality is only a conjecture). Self-solenoidization would thus be an on-line, approximate Helmholtz decomposition, as self-orthogonalization is an on-line, approximate principal component analysis; it is the learned counterpart of potential–flux decompositions of non-equilibrium dynamics [Ao, 2004, Kwon et al., 2005, Wang et al., 2008, Yan et al., 2013]. Changing only the antisymmetric component may allow new action policies without altering the perceptual model (Section 4.5).
- **Continual learning** **(HYPOTHESIS)** initial demonstration **(PUBLISHED)**. Two processes support it. First, stochastic multistability causes spontaneous replay of stored patterns, especially but not only in sequence dynamics, which consolidates old memories, as hippocampal and human replay is thought to do [Wilson and McNaughton, 1994, Liu et al., 2019] and as generative replay does in machine learning [Shin et al., 2017, van de Ven et al., 2020]. Second, self-orthogonalization stores new information only as its residual relative to what is known, so old knowledge is preserved as long as it remains at least mildly useful. This parallels orthogonal-subspace methods in machine learning [Zeng et al., 2019] and orthogonal task representations in brains and networks [Libby and Buschman, 2021, Flesch et al., 2022], as well as attractor scaffolds in hippocampal models whose memory degrades gracefully rather than collapsing [Chandra et al., 2025]. Together they give resistance to catastrophic forgetting [McCloskey and Cohen, 1989, French, 1999] without task labels or explicit regularization [Kirkpatrick et al., 2017]. The same two processes should matter on long-horizon tasks, where a system has to keep learning while it acts, without eroding what it learned earlier in the task (Section 1.1; P23).

4.4 Noise as information

In these networks noise is not a nuisance to be suppressed: its size encodes the precision of what a unit represents **(IN PREPARATION)** and partly already **(PUBLISHED)**.

Each unit’s fluctuations sample its posterior, so their variance is the posterior variance. By the fluctuation–dissipation relation of the exponential family [Kubo, 1966], the same quantity is the unit’s susceptibility, or gain. A unit is therefore most responsive where it is most uncertain, and least responsive where it is saturated and certain. Uncertainty is physically available to the system twice: as the size of its fluctuations and as its gain. This is the fluctuation–dissipation theorem put to work for computation.

In information-geometric terms, this quantity is the Fisher metric [Amari, 1998, 2016]. It is the metric in which prediction-error tangents are measured, so precision weighting is simply the geometry of that tangent space. At the network level the relevant metric is shaped by learning, through the couplings, and lowering indices with it cleanly separates the gradient and rotational parts of the dynamics (Section 4.5). In the minimal picture, the metric supplies the precision weighting that predictive coding requires [Feldman and Friston, 2010], without a separate mechanism for estimating precision: the inverse of a unit’s predictive variance is the precision of its prediction. Finite samples from a distribution produce exactly this kind of noise, so a system that learns from finite experience represents its uncertainty without extra machinery.

This view connects to sampling-based accounts in which neural variability represents uncertainty [Hoyer and Hyvärinen, 2003, Fiser et al., 2010, Buesing et al., 2011, Orbán et al., 2016], to evidence that gain variability tracks uncertainty in visual cortex [Hénaff et al., 2020], and to work on the sources and benefits of noise in neural systems [Faisal et al., 2008, McDonnell and Ward, 2011]. It is also what makes the framework a natural fit for stochastic and thermodynamic computing, which treat physical noise as a resource rather than an error [Camsari et al., 2019, Coles et al., 2023, Melanson et al., 2025].

4.5 Multistability, metastability and intrinsic action

Multistability comes from the attractors; metastability comes from what the network learns about the transitions between them.

Multistability (PUBLISHED). Multiple attractors make the networks multistable, as theory, simulations and our brain models consistently show. Noise is what makes this multistability dynamic: the system visits several attractors over time instead of settling in one, as has been proposed for resting-state activity [Deco et al., 2010, Deco and Jirsa, 2012]. A deformable attractor manifold has recently been fitted to resting-state fMRI in the same spirit [Athanasiadis et al., 2025].

Metastability (HYPOTHESIS) observed in simulations. Sequence learning produces asymmetric couplings and a non-equilibrium steady state, with transient dwelling near attractors and directed transitions between them [Kelso, 1995, Rabinovich et al., 2008, Tognoli and Kelso, 2014, Hancock et al., 2024]. Because of self-orthogonalization, we expect these dynamics to be dominated by approximately solenoidal flow over the scalar free-energy landscape (self-solenoidization, Section 4.3). Broken detailed balance in human brain dynamics [Lynn et al., 2021, Kringelbach et al., 2024] is consistent with this picture, and the non-equilibrium thermodynamics of networks with asymmetric couplings can be characterized analytically in kinetic Ising models [Aguilera et al., 2023]. In cortex, expectation modulates metastable attractor dynamics and speeds up coding [Mazzucato et al., 2019], close to the expectation-as-prior account tested in P6.

Intrinsic action: paths over a learned world model (HYPOTHESIS). The intuition is that the network learns paths over a world model it has already learned. Variational free energy (VFE) describes perception on the scalar landscape. A related quantity, possibly the expected free energy (EFE) of active inference [Friston et al., 2015, Parr et al., 2022], would describe intrinsic action, carried by the antisymmetric couplings between pairs of nodes. These couplings induce quasi-solenoidal dynamics: paths along isocontours of the landscape that determine the transitions. Without input, the same dynamics would produce spontaneous fluctuations at rest, or something like dreaming [Hobson and Friston, 2012], as in the spontaneous replay of Spisak and Friston [2026]. Whether the standard EFE arguments carry over to these networks is an open question [Millidge et al., 2021].

Decisions and planning (HYPOTHESIS). Learned transition paths could support decisions based on learned data, and planning as inference [Attias, 2003, Botvinick and Toussaint, 2012], especially across multiple timescales [Kiebel et al., 2008], in the spirit of predictive maps and replay-based planning [Pfeiffer and Foster, 2013, Stachenfeld et al., 2017]. This has a continual-learning side too: the perceptual model, or world model, need not be relearned for every new task; only the transition paths across its attractors change.

A local sketch. The mean-field dynamics and their linearization around a fixed point make the split between gradient and rotation explicit:

$$\dot{\theta}_i = -\gamma \left[\theta_i - b_i - \sum_j J_{ij} L_j(\theta_j) \right], \quad \delta \dot{\boldsymbol{\theta}} = M \delta \boldsymbol{\theta}, \quad M = -\gamma (I - JD), \quad D = \text{diag}(\psi_1''(\theta_1), \dots, \psi_n''(\theta_n)) \quad (7)$$

Here D holds each unit’s susceptibility, which is also its variance and its Fisher information (Section 4.4). Splitting the couplings into symmetric and antisymmetric parts, $J = S + A$, and lowering an index with the Fisher metric separates a symmetric, gradient-like term from an antisymmetric, rotational one:

$$DM = -\gamma(D - DSD) + \gamma DAD \quad (8)$$

For two units with $J_{12} = S + A$ and $J_{21} = S - A$, the eigenvalues are $-\gamma \pm \gamma\sqrt{(S^2 - A^2) \psi_1'' \psi_2''}$. When $|A| > |S|$, the local dynamics contain a damped rotation with frequency $\omega = \gamma\sqrt{(A^2 - S^2) \psi_1'' \psi_2''}$, which links asymmetry directly to oscillation. Symmetric couplings can admit a reversible, equilibrium model; antisymmetric couplings generically obstruct that interpretation and support directed probability currents. The cautious conclusion: asymmetry is the minimal coupling structure that distinguishes temporal consequence from static correlation, and turns an attractor landscape into directed sequence dynamics.

Directed connectivity through the attractors [\(HYPOTHESIS\)](#). In the dual picture (Section 4.3), attractors act as hidden nodes, and learned sequences live in directed, antisymmetric couplings between them. If region-to-region asymmetries serve mainly the transitions between attractors, the antisymmetric part of the couplings is low-rank and confined to the attractor subspace. With the m attractor patterns as the columns of Ξ , the $N(N - 1)/2$ free asymmetries between region pairs reduce to a small matrix Ω of attractor-to-attractor flows, with about $m(m - 1)/2$ free entries. With Λ holding the precision (curvature) of each basin, a linear-Gaussian sketch with isotropic noise gives the condition under which this flow is solenoidal, that is, leaves the stationary density of the landscape unchanged [\[Kwon et al., 2005\]](#):

$$H = \Xi \Lambda \Xi^\top, \quad T = \Xi \Omega \Xi^\top, \quad \text{solenoidal} \iff \Omega = Q \Lambda, \quad Q^\top = -Q \quad (9)$$

Exact solenoidality therefore does not require attractors of equal depth, which the present models do not have. It requires each directed coupling to be weighted by the precision of the attractor it targets, $\Omega_{\mu\nu} = Q_{\mu\nu} \lambda_\nu$: an evidence weighting of the transitions. Besides their antisymmetric part, the transition couplings then carry a small symmetric part, $\frac{1}{2} Q_{\mu\nu} (\lambda_\nu - \lambda_\mu)$, which grows with the difference in depth. Whether the predictive-coding learning rule produces this weighting by itself, since it subtracts precision-weighted predictions, is open, and simulations can measure it directly. A second route appears at the level of metastable hopping between attractors. There, added probability currents leave the occupancies unchanged exactly when they decompose into closed cycles, each added transition rate being scaled by the inverse occupancy of its source [\[Schnakenberg, 1976\]](#). A directed flow between just two attractors cannot preserve occupancies; a cycle through three or more can, whatever their depths. Both routes lead to the same testable signature: attractor occupancies set by the symmetric landscape alone, while the asymmetric couplings produce currents that leave those occupancies unchanged. Finding this despite unequal depths would itself be evidence of approximate self-solenoidization (P9).

An opportunity for directed connectivity from fMRI [\(HYPOTHESIS\)](#). Estimating directed connectivity between single pairs of regions from fMRI is notoriously hard, whether with dynamic causal modelling, Granger causality or related methods [\[Friston et al., 2003, Smith et al., 2011, Seth et al., 2015, Gilson et al., 2016\]](#), and the antisymmetric component is typically weak compared with the symmetric one. If asymmetries serve attractor transitions, estimating them at the attractor level pools evidence from many region pairs into a few attractor-to-attractor flows, which should improve statistical power; where needed, the flows can be projected back onto region pairs as $\Xi \Omega \Xi^\top$ (P22). This applies the logic of the sufficiency test (Section 5.2) to directed rather than undirected structure, and it fits the coarse-grained state spaces in which broken detailed balance has been detected in human brain dynamics [\[Lynn et al., 2021\]](#). Haemodynamic delays that differ between regions can create spurious asymmetries, and they do not average out at the attractor level; deconvolution, simulations with region-specific haemodynamic responses, and sequence tasks with a known order serve as controls.

4.6 Status

The network framework is published [\[Spisak and Friston, 2026\]](#) with open simulation code, including demonstrations of self-orthogonalization, generalization, sequence learning and resistance to catastrophic forgetting. Applied to human fMRI, connectome-based attractor networks reproduce canonical resting-state networks as approximately orthogonal attractors, beyond strict null models and replicated across datasets [\[Englert et al., 2026\]](#). Independently, a connectome-based Hopfield-type model of mouse cortex predicts fMRI co-activation patterns from its attractor landscape [\[Fasoli et al., 2026\]](#). The unit-level derivation (Section 3) is in preparation.

5. Attractors as levels of description

From the observer to the system. In Section 4.2, coarse-graining is done by an observer. Attractors are, in effect, how a system coarse-grains itself: they create slow, macroscopic variables from within, and so create the levels that a camera can later observe.

5.1 Attractors are not epiphenomena

Attractors are often dismissed as epiphenomena; in this programme they are how new timescales, and with them new levels of description, come into being (HYPOTHESIS) builds on (PUBLISHED) results.

A common objection to attractor accounts of brain function is that attractors are convenient summaries of the dynamics, while the causal work is done by neurons and synapses alone. Pressed to its end, the objection becomes a philosophical question about causation across levels of description [Kim, 1998], which we do not claim to settle. The programme does, however, offer several concrete reasons for the opposite view, each of which can be tested:

- **Attractors create timescales.** An attractor is where the dynamics slow down: the system typically dwells near it far longer than any of its units takes to relax. It is therefore a new, slower variable that the system creates for itself, and this is exactly the kind of timescale separation from which Section 3 derives perception and learning. The state of an attractor network can then act as a slowly varying field, or prior, for faster processes, so that a construction of the same kind can repeat one level up (A2, Section 8.1). Attractors are thus not a by-product of inference and learning; they are how a system produces the separation of timescales on which its inference and learning rest. Hierarchies of intrinsic timescales across the cortex are consistent with this view [Kiebel et al., 2008, Murray et al., 2014], as is the self-organization of functional hierarchies in recurrent networks with multiple timescales [Yamashita and Tani, 2008].
- **Attractors do computational work.** They can store priors and separate latent causes through self-orthogonalization; they may bridge delays in trace conditioning and working memory; and, through their transitions, they can carry sequences and intrinsic action (Sections 4.3–4.5 and 6). With symmetric couplings and approximately orthogonal attractors, the network is statistically equivalent to one in which the attractors are explicit hidden nodes that pool evidence and feed it back [Barra et al., 2012] (see Section 4.3). A description in which attractors are variables is then not a metaphor but an equivalent model.
- **Attractors make a difference under intervention.** On the interventionist account of causation [Woodward, 2004], a variable is a cause if intervening on it changes what follows. Perturbation experiments reveal attractor dynamics in exactly this sense. An activity bump imposed optogenetically on the fly’s heading circuit is maintained by the circuit [Kim et al., 2017]. Persistent activity in mouse frontal cortex returns to discrete attractor states after transient perturbations, or occasionally switches to the other attractor, followed by a switch in the animal’s action [Inagaki et al., 2019], and attractor dynamics gate the influence of distracting input on decisions [Finkelstein et al., 2021]. The programme’s own predictions are stated at the attractor level (extinction adds an attractor, chronic pain deepens one, expectations shift the posterior between them; P5, P6, P8) and can be probed with current neuroscience techniques, e.g. in closed loop with real-time fMRI (B6).
- **Attractors can form closed levels.** A macroscopic level is self-contained if it predicts its own future as well as the full microscopic state does (informational closure), and if interventions on macroscopic variables suffice to determine macroscopic outcomes (causal closure). Rosas et al. [2024] formalize such levels, which they liken to software running on physical hardware, and show that in a Hopfield network the overlaps with the stored patterns form a closed level of this kind: memory retrieval is computed at the level of the attractors, independently of microscopic detail [Hoel et al., 2013, Rosas et al., 2020]. Timescale separation is what makes slow variables approximately closed: in synergetics, slow order parameters are generated by the fast variables and in turn enslave them, a circular causality in which the macro level is anything but idle [Haken, 1983, Kelso, 1995]. Section 4.2 adds that the inferential form survives coarse-graining, so a closed level can itself infer and learn. Whether attractor-level descriptions of brain dynamics are closed in this sense is an empirical question (P17). If it is, closure has a practical corollary: for whatever is computed at that level, the attractor state is a sufficient statistic, which turns closure into a test of what is computed in the attractors (Section 5.2; P21).

Levels that can be studied in their own terms. Simon [1962] argued that complex systems evolve far more readily when they can build on stable intermediate forms, and that the result is a nearly decomposable hierarchy, in which fast interactions within subsystems are separated from slow, aggregate interactions between them. Attractors are dynamical counterparts of such stable intermediate forms. Where levels are approximately closed in this sense, each can be studied largely in its own terms, without simulating everything beneath it, and an elementary step that recurs at each level needs to be understood only once (Section 11).

5.2 What is computed in the attractors? Counterfactuals and a sufficiency test

For brain data, the programme aims at mechanism, not at predictive accuracy. Two kinds of test follow: counterfactual simulation, and a test, derived from closure, of which quantities are carried by large-scale attractors and which are computed below them (HYPOTHESIS) proposed method.

Mechanism, not accuracy. Data-driven predictive models of brain and behaviour are optimized to predict, and a mechanistic model need not beat them at it. Its value lies in answering questions they cannot answer: what would happen under an intervention that is absent from the training data. The fitted networks are generative and can be manipulated in silico: an attractor can be deepened or flattened, prior precision raised or lowered, the biases of selected regions shifted, or couplings altered. The decisive test is whether such counterfactual simulations, specified in advance, predict the effects of real manipulations, such as placebo and nocebo instructions, conditioning and extinction, pharmacological challenge, non-invasive brain stimulation or closed-loop real-time fMRI. The comparison is with mechanistically different variants of the same model (couplings fixed, biases only, a landscape without attractor structure), not with black-box predictors. Conventional predictive models keep an important role, however: as shown next, they set the reference that an attractor-level description has to reach, not beat.

Two levels in one fit. Score matching separates two components of the regional dynamics (graphical abstract, Line B). The couplings define the landscape and its attractors: dynamics at the network level. The biases capture what drives each region from below and from outside: mesoscale computation within regions, and sensory input. Both can be read out as time series from the same data: the attractor state $a(t)$, given by the overlaps of the regional state $x(t)$ with the fitted attractors (or by the attractor towards which each state relaxes), and the input $b(t)$ inferred by the fitted model.

Closure implies sufficiency. If the attractor level is closed in the sense of Section 5.1 (P17), and a quantity y (a perceptual report, a behavioural choice, an experimental condition or a clinical variable) is computed at that level, then the regional data carry no information about y beyond what the attractor state carries. The attractor state is then a sufficient statistic for y :

$$p(y | x) = p(y | a(x)) \iff I(y; x) = I(y; a) \quad (10)$$

Operationally, y should be predictable from the attractor time series just as well as from the full regional data; not better, but without loss. Because $a(t)$ and $b(t)$ are both computed from the regional time series by the fitted model, the chain rule of mutual information splits what the data carry about y exactly into three parts:

$$I(y; x) = \underbrace{I(y; a)}_{\text{attractor level}} + \underbrace{I(y; b | a)}_{\text{upstream input}} + \underbrace{I(y; x | a, b)}_{\text{neither}} \quad (11)$$

Quantities for which the first term exhausts the total are computed in attractor space. Those that need the second term are carried by upstream, mesoscale processing or by sensory drive. A substantial third term points to structure that the model misses. In this way closure, a property of a level, becomes a criterion for what lives on that level.

Mixed cases and controls. Many brain functions will be mixed, and the attractor share $I(y; a)/I(y; x)$ is then a graded measure rather than a verdict. Its value is that it is principled rather than descriptive: the criterion follows from the theory, not from choosing whichever feature set predicts best. Three controls keep it honest. First, sufficiency must be specific: equally compact coarse-grainings that ignore the attractors (principal components, random projections, canonical networks of the same dimension) should lose information that the attractor coordinates keep; otherwise the result reflects compression in general rather than attractors. Second, the information terms are estimated from the cross-validated performance of nested models fitted

with the same learner, and are bounded in practice by sample size, model capacity and haemodynamic blurring; equal performance is therefore established with equivalence tests, not read off a non-significant difference. Third, sufficiency for an external quantity is necessary but not sufficient for closure, which also requires the attractor level to predict its own future (P17); the two tests are run together. The same logic applies between resolutions: replacing x with data from a finer atlas, or at voxel level, asks whether the coarse attractors suffice for what the finer data carry (Section 4.2).

An example: pain. Pain is a natural test case because it combines strong input with strong expectations. If expectations act as attractor priors (P6), placebo and nocebo effects on reported pain should be largely attractor-sufficient, and an in-silico shift of the corresponding prior should reproduce them. The effect of stimulus intensity, by contrast, should be carried mostly by the biases, as nociceptive drive. In chronic pain the same decomposition can be followed over time, to see whether chronification moves pain from an input-carried to an attractor-level phenomenon, or the reverse (P8).

6. Emergent, biology-like phenomena: hypotheses

Several well-known features of biological learning may follow from the same few properties, above all self-orthogonalization. Below we highlight a few selected examples; further-reaching possibilities are collected in Section 11. Each is a hypothesis to be tested, not a result.

- **Primacy effect** (HYPOTHESIS). Because new information is stored only as a residual relative to existing attractors, early material serves as the reference frame for what follows and should be remembered better. This would mirror the primacy effect in human memory [Murdock, 1962], in contrast to classical palimpsest networks, which forget their oldest memories first [Nadal et al., 1986, Parisi, 1986].
- **Extinction, renewal and reinstatement** (HYPOTHESIS). Self-orthogonalization restricts how representations can change:
 - *Acquisition.* A newly learned representation is stored in an attractor approximately orthogonal to all previous ones, so only its residual is stored, and old representations shift as well (remapping) [Colgin et al., 2008].
 - *Context.* Overlap with old memories, for example an association with fear or pain, induces correlations that must be balanced by non-overlapping components to keep the attractors orthogonal. This is how context A becomes part of the memory.
 - *Extinction.* The old memory could only be erased if everything else stayed unchanged, which is practically impossible. Instead a new attractor forms, again orthogonal to the others. Because the stimulus is similar or identical but the outcome differs, self-orthogonalization puts strong pressure on using context to tell the two situations apart: context B.
 - *Renewal and reinstatement.* The old context, or the unconditioned stimulus itself, therefore brings back the original response; through sequence learning, the network may even anticipate it.
 - *Related work.* This mechanistic account parallels the evidence that extinction is new, context-dependent learning [Rescorla and Heth, 1975, Bouton, 2004, Rescorla, 2004], latent-cause models in which extinction creates a new latent cause [Gershman et al., 2010, 2017], state-splitting accounts of renewal [Redish et al., 2007] and the neurobiology of fear extinction [Milad and Quirk, 2012].
- **Trace conditioning and working memory** (HYPOTHESIS). Self-sustaining attractor activity that persists after the stimulus has gone would bridge the interval in trace conditioning [Clark and Squire, 1998] and support working memory, as in classical persistent-activity models [Goldman-Rakic, 1995, Amit, 1997, Wang, 2001]. Activity-silent accounts [Stokes, 2015] provide the contrast to test against.
- **Oscillations during acquisition** (HYPOTHESIS). When a new stimulus is not well explained by any existing attractor, ghost attractor or superposition, only oscillatory sampling across attractors remains. This is slow and costly, because the boundaries between attractors have high free energy, so learning will try to store the new information in an attractor, or at least a ghost, while keeping orthogonality (remapping). The oscillations should then fade as learning proceeds. This may relate to theta-band activity during acquisition and encoding [Buzsáki, 2002, Hasselmo et al., 2002, Seidenbecher et al., 2003, Likhtik et al., 2013], and to the oscillations that emerge in circuits optimized for sampling-based inference [Aitchison and Lengyel, 2016, Echeveste et al., 2020].

- **Curiosity and exploration** (HYPOTHESIS) open. How far the usual active-inference arguments for epistemic value [Friston et al., 2015, 2017] carry over to these networks is unclear, but uncertainty here naturally encourages exploration [Oudeyer, 2007, Schmidhuber, 2010, Gottlieb and Oudeyer, 2018]. In this picture a “dark room” [Friston and Ao, 2012] is not predictable: it is a lack of evidence, which leaves spontaneous activity to sample the priors and may produce hallucination-like states. This links to accounts of hallucinations as strong priors [Powers et al., 2017, Corlett et al., 2019] and to pain without input, as in phantom-limb pain [Flor et al., 2006].
- **Reinforcement learning** (HYPOTHESIS). The networks are compatible with reinforcement learning in the way active-inference accounts propose [Friston et al., 2009, Sutton and Barto, 2018], but this relies on fixed priors set by evolution, such as pain. Pain research is therefore a good test bed [Büchel et al., 2014, Wager and Atlas, 2015, Tabor et al., 2017].

7. Predictions and falsification criteria

Every prediction below is stated in advance, with the test bed and the result that would count against it. Self-orthogonalization is the system’s most prominent signature, but it is hard to test in high dimensions: random vectors also tend towards orthogonality. Real inputs are almost always structured and correlated, so orthogonality of learned attractors is not trivial. Still, the decisive test we propose is the projector-network signature (P2), not orthogonality alone. All geometric tests use strong null models: phase-randomized time series, symmetry-preserving permutations, covariance-matched Gaussian models, and condition-label shuffling. Analysis plans will be preregistered where possible. For the neuroscience predictions the criterion is mechanistic: counterfactual simulations specified in advance, and the sufficiency test of Section 5.2, rather than predictive accuracy beyond that of existing data-driven models.

#	Prediction	Test bed	Would count against it
P1	Learned attractors are approximately orthogonal, even for correlated inputs	Rest and task fMRI; simulations	Attractor angles indistinguishable from matched null models
P2	Decisive test of self-orthogonalization: the networks approach a Kanter-Sompolinsky projector network, whose attractors coincide with the leading eigenvectors of the coupling matrix	fMRI-fitted networks; simulations	Eigenvector-attractor alignment no better than null models
P3	Orthogonalization depends on the timing of the learning rule: pairing the error with the renewed state orthogonalizes; pairing it with the preceding state merges attractors	Simulations; physical hardware	Orthogonalization with reversed pairing, or none with the correct pairing
P4	Noise encodes precision: a unit's variance equals its gain, and variability tracks inferred precision	Physical hardware; neural and BOLD variability vs responsiveness	No relation between variability and gain or precision
P5	Extinction creates a new, approximately orthogonal attractor rather than erasing the old one; renewal is a context-driven return to it	Threat-learning and extinction fMRI; benchmark battery	Extinction as weakening of the original attractor; renewal not predicted by the fitted landscape
P6	Expectations act as attractor priors: placebo and nocebo shift the posterior towards expectation-congruent states, scaling with prior precision	Placebo and nocebo fMRI	An in-silico shift of the prior does not reproduce the observed placebo and nocebo effects, or these effects are not attractor-sufficient (P21)
P7	Task activity decomposes into network-level (couplings) and input-driven (biases) components	Task fMRI; 7T laminar fMRI (feedback vs feed-forward layers)	Decomposition does not generalize across tasks or maps onto the wrong layers
P8	Chronic pain reflects a deepened, harder-to-leave pain attractor (macro-scale) and/or region-intrinsic change (meso-scale), distinguishable within subjects over time by the attractor-input decomposition (P21)	Longitudinal phantom-limb pain cohort	Neither component tracks chronification
P9	Self-solenoidization: sequential experience produces asymmetric couplings whose flow is approximately solenoidal (along free-energy isocontours) ²⁰ and replays sequences in their learned direction, while attractor occupancies stay those	Sequence-learning fMRI (Helmholtz decomposition of fitted flows); simulations	Transition flow mainly along the free-energy gradient, no directed replay, or occupancies shifted by the asymmetric couplings

Negative results are informative by design: they delimit where the construction applies and point to mechanisms missing from the theory, and they will be published.

8. Technical plan

The three lines share one model class, one set of definitions and one set of predictions; results from each are checked against the others and are often published together.

8.1 Line A – Approach and theory

- **A1. Complete and test the unit-level construction.** Publish the derivation; test its contingent predictions in finite jump-process simulations in which no inference or learning rule is coded: the form of the nonlinearity under different resting laws, fluctuation–gain equality, and the dependence of orthogonalization on learning-rule timing (P3, P4).
- **A2. Composition and hierarchy.** Formalize how evidence at one level acts as the effective field at the next, and how nested timescales give rise to hierarchical message passing. Derive principled coarse-graining criteria: bounds on how well attractors are preserved as a function of within-parcel heterogeneity, and conditions under which the coarse transition law, not only the attractor set, is preserved (Section 4.2). Show how attractors at one level supply the slow timescale for the next, and characterize when the resulting levels are informationally and causally closed (Section 5.1; P17).
- **A3. Action and planning.** Develop the network-level account of action: solenoidal flow from antisymmetric couplings as intrinsic action, and effector-driving nodes as extrinsic action. Study planning as sequence prediction over the learned landscape, and the link between learnable surprise and exploration.
- **A4. Thermodynamic cost.** Relate the quantities of the construction (alignment gap, learnable versus unlearnable surprise) to the heat dissipated during perception and learning, as a bridge to physical implementation. Because each synapse is the thermodynamic Legendre conjugate of a correlation, learning has a thermodynamic price that can be computed in theory and measured in hardware.

Methods: non-parametric information geometry, stochastic thermodynamics, averaging and adiabatic theory, and large-scale simulation.

8.2 Line B – Brains: generative models of large-scale brain dynamics

- **B1. Model fitting.** With the open-source package *attractome*, fit networks to resting-state, task and naturalistic fMRI by score matching with haemodynamic deconvolution, estimating couplings (network-level dynamics) and biases (upstream, mesoscale input). Validate generatively, by simulating held-out data, and against null models. This includes individual-level models from dense, repeatedly sampled single-subject datasets. Before interpretation, the fitted landscapes are tested across resolutions: project and match across nested atlases, measure homogeneity, and compare with random parcellations, in population and single-subject data (Section 4.2; P18–P20).
- **B2. Tests of the theory.** Attractor orthogonality and eigenmode alignment (P1, P2), the couplings–input decomposition (P7), self-solenoidization, transition learning and directed attractor-to-attractor flows in sequence tasks (P9, P13, P22), prior mixing and oscillatory alternation between attractors (P10, P15), spontaneous sampling of priors at rest (P16), and the closure of the attractor-level description (P17) and its corollary, attractor sufficiency (P21). Ultra-high-field (7T) laminar fMRI tests whether inferred priors and inputs map onto feedback and feedforward cortical layers.
- **B3. Learning without forgetting: extinction and renewal (P5).** Fit landscapes before and after acquisition, extinction and renewal in threat-learning paradigms, and test whether extinction adds a new, approximately orthogonal attractor.
- **B4. Expectations as priors (P6).** Model placebo and nocebo effects on pain as precision-weighted combinations of attractor priors and nociceptive input, using placebo-imaging consortium data. Test whether in-silico shifts of the prior, specified in advance, reproduce the observed effects, and whether placebo responses are attractor-sufficient while the effects of stimulus intensity are carried by the input (Section 5.2).

- **B5. Plasticity in pain chronification (P8).** Separate meso-scale (region-intrinsic) from macro-scale (network) plasticity in chronic low back pain, fibromyalgia and longitudinal phantom-limb pain cohorts, using the attractor–input decomposition to follow, within subjects, the level at which pain is carried as it becomes chronic (P8, P21).
- **B6. From predictive to mechanistic neuromarkers.** Brain-based predictive models of pain, cognition and clinical outcomes have matured: they can be replicable with moderate sample sizes [Spisak et al., 2023], and some generalize to independent patient populations [Spisak et al., 2020]. Most, however, rest on data-agnostic machine learning, so their behaviour out of context is hard to predict, their features are hard to map onto mechanisms, and they cannot simulate what has not been measured. The aim here is not to outperform them but to explain what they capture. The sufficiency test (Section 5.2) asks whether the information a predictive model exploits is carried by large-scale attractors or by upstream input, with the predictive model itself as the reference to be matched. Compact model-based features (attractor depth, dwell times, transition costs) are validated externally across cohorts, and counterfactual simulation generates intervention hypotheses that are specified in advance. Real-time fMRI with reinforcement learning from brain feedback provides closed-loop, causal tests of these hypotheses.

Data: large open datasets and cohorts from established collaborations (placebo and pain consortia, two DFG collaborative research centres, an ERA-NET pain consortium).

8.3 Line C – Computers: emergent learning as an alternative route to AI

- **C1. Software and scaling.** Efficient, open implementations; map the networks’ operational regimes (temperature, noise, input precision) and scaling behaviour.
- **C2. A neuroscience-inspired benchmark battery.** Continual learning, extinction, reinstatement and renewal, interference, primacy, reversal and sequence tasks, matched to the paradigms of Line B so that brain and machine are compared on the same tasks (P5, P11, P14), and long-horizon tasks scored against a repeated-sampling baseline (P23). Baselines: standard continual-learning methods (regularization- and replay-based) and backpropagation-trained recurrent networks.
- **C3. Uncertainty and generative inference.** Bayesian reconstruction and denoising as a function of prior and input precision; calibration of the networks’ uncertainty; generalization through prior mixing (P10).
- **C4. Native physical implementation (P12).** Stochastic binary units are the tanh case of the minimal unit, which makes probabilistic-bit and related hardware a natural fit. Steps: simulate hardware-constrained networks; run inference natively on hardware with the local learning rule applied from the measured states; then co-design slowly adapting couplings native to the substrate. Candidate substrates: thermodynamic, memristive, nanomagnetic and related emerging architectures [Camsari et al., 2019, Coles et al., 2023, Harabi et al., 2022], with partners. Device evaluation has so far relied on emulation metrics that rate how faithfully a device reproduces biological synapses, neurons or networks [Bruno et al., 2026]. As a complement, the construction suggests theory-derived figures of merit: the separation of timescales between state and conductance, a bounded, maximum-entropy response, noise variance that tracks gain, bidirectional local updates with the prescribed timing, support for asymmetric couplings, and, at the network level, attractor orthogonality and resistance to forgetting.
- **C5. Closed-loop agency.** Maze, reversal and detour tasks probe world-model formation, planning and autonomous action, and bridge back to behavioural neuroscience.

9. Risks, limitations and fallbacks

The programme is designed so that each risk, if it materializes, still produces an informative result.

Risk	Consequence	Fallback or what is learned
Assumptions of the construction do not hold for a given substrate	The inferential reading is not guaranteed there	Delimits the scope of the theory; the network models remain usable as phenomenological generative models
Orthogonality in brain data is explained by high-dimensional chance or measurement artefacts fMRI is too slow or too noisy for some signatures	P1 and P2 fail Weak tests in Line B	Strong null models detect this; attention shifts to dynamical predictions (P5–P9) Laminar 7T fMRI, dense single-subject sampling, and simulations with matched noise; electrophysiology via collaborators
Counterfactual simulations fail to predict the effects of real manipulations	The fitted landscapes describe brain dynamics without explaining them	Model variants with fixed couplings or biases locate which component fails; the sufficiency analysis still maps which quantities are carried at the attractor level; results are published either way
Networks do not scale on conventional hardware	Line C limited to small and medium problems	Focus on properties (continual learning, calibrated uncertainty) rather than scale; emerging hardware as the long-run route
Emerging hardware cannot realize slow adaptive couplings	Native learning (C4) not reached	Hardware-in-the-loop learning remains publishable; results inform hardware development
Brain and machine share no common signatures	The common-principle hypothesis fails	A clear negative result: the theory may describe a viable route to learning that brains do not use
Attractor-level descriptions of brain data are not closed (P17 fails)	Attractors may be better treated as descriptions than as a level in their own right	The fitted networks remain generative models of regional dynamics; closure is then tested at finer scales and in simulations, where timescale separation is larger (A2)
Few quantities are attractor-sufficient; most are mixed (P21)	No clean assignment of functions to levels	The attractor share remains a graded, theory-derived measure; a consistently large input share would itself show that large-scale attractors carry less of the computation than proposed
Attractors do not match across resolutions (P18 or P19 fails)	fMRI-fitted networks cannot be read as views of one landscape	Claims are restricted to resolutions validated by held-out generative tests; precision fMRI, 7T data and simulations with known ground truth locate where homogeneity fails

Limitations and open questions

Much of this programme rests on heuristics that still have to be made rigorous; we list the main ones so that readers can weigh them.

- **Sufficient, not necessary.** The unit-level construction shows one route from physics to inference and learning. It does not show that brains, or any particular material, take this route.
- **Strong assumptions.** Timescale separation, bounded coordinates and leading-order interactions are idealizations. Where they fail, the inferential reading is not guaranteed.

- **Heuristics not yet derived.** Self-solenoidization, the duality with antisymmetric couplings, prior mixing through oscillatory alternation, and the role of expected free energy in intrinsic action are hypotheses, not results.
- **Not an exact gradient.** The network learning rule is not the exact gradient of a single objective, so its stationary behaviour is characterized only approximately.
- **Minimal units.** One-dimensional, continuous-Bernoulli units are the simplest case; multidimensional or spiking units need extensions.
- **Measurement.** fMRI is slow, indirect and confounded by vascular and physiological signals. Results depend on coarse-graining choices, since a parcellation shows only distinctions that live across parcels (Section 4.2), and orthogonality in high dimensions needs strong null models.
- **Scale.** Demonstrations so far are small to medium in scale. We make no claim of competitiveness with large deep-learning systems.
- **Speculative links.** The broader view of life as stacked timescales and the connections to quantum-like geometry and to consciousness are collected in the Outlook (Section 11); the rest of the programme does not depend on them.
- **The roots, not the whole tree.** The programme aims to explain where the brain’s capacities come from, not the full complexity with which a particular brain realizes them (Section 1.3). Success would leave most of neuroscience’s hard problems open, while giving them, we hope, a common starting point.
- **Levels and causation.** Whether attractors are causes or descriptions ultimately touches philosophical questions that the programme does not settle [Kim, 1998]. We argue only for causal relevance in senses that can be tested: the creation of timescales, the effects of interventions, and closure (Section 5.1; P17).

10. Relation to other approaches

The programme complements, rather than competes with, neighbouring approaches; each meets it at a different point.

Approach	Central object	Where it meets this programme	What this programme adds
Neuroconnectionism, task-optimized deep networks	Architecture, objective and training data that make networks brain-like	Shared aim: computational models tested against brain and behavioural data	A first-principles route to the learning rule and dynamics; whole-brain generative modelling; continual learning by construction
Biophysical whole-brain models (neural mass, neural field)	Biophysics at one chosen scale	Shared target: large-scale brain dynamics	A scale-agnostic family of rules that complements scale-specific mechanisms; task fidelity, as networks trainable as learners
Data-driven reconstruction of brain dynamics (recurrent networks fitted to recordings)	Generative dynamical models learned from data	Shared goal: generative models of brain dynamics	A model class derived from first principles, with interpretable priors, inputs and learning
Free Energy Principle, active inference	Variational free energy minimization by self-organizing systems	The networks minimize variational free energy at two timescales	Sufficient physical conditions for one minimal system; no Markov blanket required; falsifiable predictions
Bayesian brain, predictive coding	Posterior inference; prediction errors	Prediction errors, posteriors and precisions appear in the construction	A physical origin for prediction errors and precision; the timing of the learning rule
Sampling-based probabilistic computation	Neural variability as samples from a posterior	Fluctuations sample the posterior at two timescales	Variance equals gain; noise as precision; a learned Fisher metric
Hopfield networks, Boltzmann machines, dense associative memories	Energy landscapes; contrastive learning; storage capacity	The networks generalize stochastic Hopfield networks and Boltzmann machines	Local, single-phase learning; asymmetric couplings; self-orthogonalization
Landscape-flux theory of non-equilibrium systems	Decomposition into potential and curl flux	Symmetric and antisymmetric couplings split gradient and solenoidal flow	A learning rule expected to produce this split on-line (self-solenoidization)
Latent-cause models of learning and extinction	Inference over hidden causes of experience	Extinction as the formation of a new cause	An attractor-level mechanism with neural predictions
Physical learning (equilibrium propagation, coupled learning)	Learning rules native to physical networks	Shared aim: learning native to physics	Unsupervised, generative learning derived from first principles; a link to brain data
Continual-learning methods in machine learning	Avoiding catastrophic forgetting	Shared benchmarks	Resistance to forgetting as an emergent property, not an add-on
Information-theoretic accounts of emergence (causal emergence; informational and causal closure)	Macroscopic variables that are self-contained, or causally more informative than the micro level	Attractors as candidate closed macro-variables of brain dynamics	A mechanism that generates such levels from within, through learned attractors with their own timescales; an fMRI test (P17)
Connectome-based attractor and energy-landscape models of brain data	Energy landscapes or attractors of Ising- or Hopfield-type models fitted to, or built from, brain connectivity	The same model family and the same data [Watanabe et al., 2013, Ezaki et al., 2017, Gu et al., 2018, Fasoli et al., 2026]	A derived learning rule; attractors read as Bayesian priors; self-orthogonalization as a testable signature (P1, P2)
Evolution and development as learning	Associative learning by networks of interacting components over evolutionary and developmental time	Learning as a property of organized matter, across levels [Watson et al., 2011, Vanchurin et al., 2022, Watson and Levin, 2023, Fields and	A minimal physical construction of the elementary step, with tests in brains and machines; a reading of evolution as refining, rather than in-

Intellectual lineage and acknowledgements

This programme stands on many traditions, and most of its ideas are new combinations of theirs rather than independent innovations. We acknowledge in particular:

- **Statistical physics and information geometry:** maximum entropy [Jaynes, 1957], fluctuation–dissipation [Kubo, 1966], projection operators and memory kernels [Zwanzig, 1961, Mori, 1965], averaging and multiscale methods [200, 2008], synergetics and slaving [Haken, 1983], centre manifolds [Carr, 1981], non-parametric information geometry [Pistone and Sempi, 1995] and information geometry more broadly [Amari, 1998, 2016].
- **The free energy principle and active inference:** Friston [2010, 2019], Palacios et al. [2020], Parr et al. [2022], Friston et al. [2023a,b], Ramstead et al. [2023], Sakthivadivel [2022] and Friston et al. [2025]; and the critical literature that sharpened our formulation [Andrews, 2021, Biehl et al., 2021, Millidge et al., 2021, Aguilera et al., 2022, Bruineberg et al., 2021].
- **The Bayesian brain and predictive coding:** Rao and Ballard [1999], Knill and Pouget [2004], Friston [2005], Feldman and Friston [2010], Bastos et al. [2012], Whittington and Bogacz [2017], Song et al. [2024] and Senn et al. [2024]; and hierarchical Bayesian models of learning under uncertainty [Mathys, 2011].
- **Sampling-based probabilistic computation:** Hoyer and Hyvärinen [2003], Fiser et al. [2010], Berkes et al. [2011], Buesing et al. [2011], Kappel et al. [2015], Aitchison and Lengyel [2016], Orbán et al. [2016], Echeveste et al. [2020] and Hénaff et al. [2020]; and probabilistic population codes as an alternative [Ma et al., 2006].
- **Attractor networks and energy-based models:** Hopfield [1982], Amit et al. [1985], Ackley et al. [1985], Smolensky [1986], Personnaz et al. [1986], Sompolinsky and Kanter [1986], Kleinfeld [1986], Kanter and Sompolinsky [1987], Amit [1989], Hinton [2002], Barra et al. [2012], Krotov and Hopfield [2016], Ramsauer et al. [2020], Khona and Fiete [2022], Hoover et al. [2025] and Clark [2026]; Hebbian unlearning and dreaming [Hopfield et al., 1983, Fachechi et al., 2019, Serricchio et al., 2025]; and perturbation evidence for attractor dynamics in brains [Kim et al., 2017, Inagaki et al., 2019, Finkelstein et al., 2021].
- **Local Hebbian and anti-Hebbian learning:** Oja [1982], Földiák [1990], Bi and Poo [1998], Gerstner et al. [2018], Pehlevan and Chklovskii [2019] and Lillicrap et al. [2020].
- **Non-equilibrium dynamics and metastability:** potential–flux decompositions [Ao, 2004, Kwon et al., 2005, Wang et al., 2008, Friston and Ao, 2012, Yan et al., 2013], irreversible samplers [Ma et al., 2015, Rey-Bellet and Spiliopoulos, 2015], metastability and itinerancy [Kelso, 1995, Tsuda, 2001, Rabinovich et al., 2008, Tognoli and Kelso, 2014, Hancock et al., 2024] and broken detailed balance in the brain [Lynn et al., 2021, Kringelbach et al., 2024] and the non-equilibrium thermodynamics of asymmetric networks [Aguilera et al., 2023].
- **Large-scale brain dynamics:** resting-state networks and their relation to tasks [Raichle et al., 2001, Fox et al., 2005, Smith et al., 2009, Cole et al., 2016, Margulies et al., 2016, Dosenbach et al., 2025]; whole-brain models [Deco et al., 2008, 2010, Honey et al., 2009, Sanz Leon et al., 2013, Breakspear, 2017]; maximum-entropy and energy-landscape models [Schneidman et al., 2006, Watanabe et al., 2013, Ezaki et al., 2017, Gu et al., 2018]; connectome-based attractor models of fMRI [Englert et al., 2026, Fasoli et al., 2026]; attractor manifolds and timescale separation [Runfola et al., 2025, Athanasiadis et al., 2025]; metastable attractor dynamics and expectation [Mazzucato et al., 2019]; parcellation and its evaluation [Gordon et al., 2014, Eickhoff et al., 2018, Urchs et al., 2019] and precision functional mapping [Gordon et al., 2017]; brain-state models [Vidaurre et al., 2017]; data-driven reconstruction of brain dynamics [Koppe et al., 2019, Singh et al., 2020, Durstewitz et al., 2023]; and the debate between bottom-up and top-down neuroscience [Rosas et al., 2025].
- **Continual learning:** McCloskey and Cohen [1989], McClelland et al. [1995], French [1999], Kirkpatrick et al. [2017], Shin et al. [2017], Zeng et al. [2019], Hadsell et al. [2020], van de Ven et al. [2020], Libby and Buschman [2021], Flesch et al. [2022] and Chandra et al. [2025].
- **Learning, extinction and expectation:** Rescorla and Heth [1975], Bouton [2004], Rescorla [2004], Redish et al. [2007], Gershman et al. [2010], Milad and Quirk [2012], Büchel et al. [2014], Wager and Atlas [2015] and Tabor et al. [2017].
- **Physical, neuromorphic and thermodynamic computing:** Strukov et al. [2008], Still et al. [2012], England [2015], Prezioso et al. [2015], Goldt and Seifert [2017], Scellier and Bengio [2017],

Camsari et al. [2019], Borders et al. [2019], Conte et al. [2019], Hooker [2021], Stern et al. [2021], Schuman et al. [2022], Coles et al. [2023], Harabi et al. [2022], Stern and Murugan [2023] and Stern et al. [2025].

- **NeuroAI:** Hassabis et al. [2017], Lake et al. [2016], Richards et al. [2019], Ali et al. [2022], Doerig et al. [2023] and Zador et al. [2023].
- **Scale-free and basal cognition:** Lyon [2005], Watson et al. [2011], Levin [2019], Fields and Levin [2022], Vanchurin et al. [2022], Watson and Levin [2023], Friston et al. [2025] and Yampolskaya and Mehta [2026].
- **Emergence, levels and causation:** hierarchy and near-decomposability [Simon, 1962], “more is different” [Anderson, 1972], order parameters and circular causality [Haken, 1983, Kelso, 1995], hierarchies of timescales in the brain [Kiebel et al., 2008, Murray et al., 2014] and in recurrent networks [Yamashita and Tani, 2008], interventionist causation [Woodward, 2004], causal emergence and closure [Hoel et al., 2013, Rosas et al., 2020, 2024]; and, for the opposing view, the causal exclusion argument [Kim, 1998].

We are grateful to the many colleagues whose work and conversations shaped these ideas; any errors in combining them are our own.

11. Outlook: further-reaching possibilities

The rest of the programme does not depend on the ideas in this section **(SPECULATIVE)**.

We state them because they follow naturally from the construction, and because stating them openly makes it easier to see where it might lead, or where it might break.

Life as stacked timescales. From the point of view of Section 5, life stacks timescales, from molecules and cells to circuits, behaviour, development and evolution, and it appears to do so in a way that largely respects the boundaries between levels (Section 5.1). Hopfield-type attractor dynamics indeed recur across biology, from cell differentiation and epigenetic memory to molecular self-assembly [Yampolskaya and Mehta, 2026]; local adaptation of interactions can produce associative memory at the scale of a whole system [Watson et al., 2011]; and evolution and development have been cast as learning across nested levels and timescales [Vanchurin et al., 2022, Watson and Levin, 2023]. That living systems largely respect these boundaries is a great gift for scientific understanding, and it is what makes the programme’s bet plausible: if the same elementary step recurs at each level, understanding it once gives a handle on many.

Evolution as refining a law rather than inventing rules. If the elementary step is a property of organized matter rather than a biological invention (Sections 1.3 and 3.2), evolution may have worked less like an engineer designing rules and more like a craftsman perfecting a material. It would start from structures that align only weakly with their environment and favour substrates in which the same law operates with more headroom: more reliable separation of timescales, wider ranges of coupling, richer architectures. Because better-aligned systems can in turn build better substrates, the gains could compound, so that alignment improves at an accelerating, perhaps exponential, rate over evolutionary time. This echoes views of evolution and development as one scale-free process acting on materials that are already competent [Fields and Levin, 2020, Levin, 2023], and of evolution as learning [Watson and Levin, 2023, Vanchurin et al., 2022]. Whether any of this can be made quantitative is open.

Quantum-like geometry. The rotational, solenoidal part of the dynamics (Section 4.5) leads into the territory of tangent and cotangent bundles. On a dually flat statistical manifold, Dombrowski’s construction gives the tangent bundle a Kähler structure [Dombrowski, 1962, MOLITOR, 2012, Molitor, 2013], the same kind of geometry that underlies the geometric formulation of quantum mechanics [Kibble, 1979, Ashtekar and Schilling, 1999]. Closed solenoidal loops may then accumulate geometric, Berry-like phases [Berry, 1984], which are known to arise in classical stochastic kinetics [Sinitsyn and Nemenman, 2007]. Whether this formal kinship can account for quantum-like phenomena in cognition and decision making [Busemeyer and Bruza, 2012, Pothos and Busemeyer, 2013] is open.

Consciousness. Attractors offer a plausible latent dimension for a global workspace [Baars, 1988, Dehaene et al., 1998, Mashour et al., 2020, Whyte and Smith, 2021] or for the “inner screen” of Ramstead et al. [2023]. We make no claim here beyond pointing out these links.

References

- Springer New York, 2008. ISBN 9780387738291. doi:[10.1007/978-0-387-73829-1](https://doi.org/10.1007/978-0-387-73829-1). URL <http://dx.doi.org/10.1007/978-0-387-73829-1>.
- David H. Ackley, Geoffrey E. Hinton, and Terrence J. Sejnowski. A Learning Algorithm for Boltzmann Machines*. *Cognitive Science*, 9(1):147–169, 1 1985. ISSN 1551-6709. doi:[10.1207/s15516709cog0901_7](https://doi.org/10.1207/s15516709cog0901_7). URL http://dx.doi.org/10.1207/s15516709cog0901_7.
- Miguel Aguilera, Beren Millidge, Alexander Tschantz, and Christopher L. Buckley. How particular is the physics of the free energy principle? *Physics of Life Reviews*, 40:24–50, 3 2022. ISSN 1571-0645. doi:[10.1016/j.plrev.2021.11.001](https://doi.org/10.1016/j.plrev.2021.11.001). URL <http://dx.doi.org/10.1016/j.plrev.2021.11.001>.
- Miguel Aguilera, Masanao Igarashi, and Hideaki Shimazaki. Nonequilibrium thermodynamics of the asymmetric Sherrington-Kirkpatrick model. *Nature Communications*, 14(1), 2023. ISSN 2041-1723. doi:[10.1038/s41467-023-39107-y](https://doi.org/10.1038/s41467-023-39107-y). URL <http://dx.doi.org/10.1038/s41467-023-39107-y>.
- Laurence Aitchison and Máté Lengyel. The Hamiltonian Brain: Efficient Probabilistic Inference with Excitatory-Inhibitory Neural Circuit Dynamics. *PLOS Computational Biology*, 12(12):e1005186, 12 2016. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1005186](https://doi.org/10.1371/journal.pcbi.1005186). URL <http://dx.doi.org/10.1371/journal.pcbi.1005186>.
- Abdullahi Ali, Nasir Ahmad, Elgar de Groot, Marcel Antonius Johannes van Gerven, and Tim Christian Kietzmann. Predictive coding is a consequence of energy efficiency in recurrent neural networks. *Patterns*, 3(12):100639, 12 2022. ISSN 2666-3899. doi:[10.1016/j.patter.2022.100639](https://doi.org/10.1016/j.patter.2022.100639). URL <http://dx.doi.org/10.1016/j.patter.2022.100639>.
- Shun-ichi Amari. Natural Gradient Works Efficiently in Learning. *Neural Computation*, 10(2):251–276, 2 1998. ISSN 1530-888X. doi:[10.1162/089976698300017746](https://doi.org/10.1162/089976698300017746). URL <http://dx.doi.org/10.1162/089976698300017746>.
- Shun-ichi Amari. *Information Geometry and Its Applications*. Springer Japan, 2016. ISBN 9784431559788. doi:[10.1007/978-4-431-55978-8](https://doi.org/10.1007/978-4-431-55978-8). URL <http://dx.doi.org/10.1007/978-4-431-55978-8>.
- D. Amit. Model of global spontaneous activity and local structured activity during delay periods in the cerebral cortex. *Cerebral Cortex*, 7(3):237–252, 4 1997. ISSN 1460-2199. doi:[10.1093/cercor/7.3.237](https://doi.org/10.1093/cercor/7.3.237). URL <http://dx.doi.org/10.1093/cercor/7.3.237>.
- Daniel J. Amit. *Modeling Brain Function: The World of Attractor Neural Networks*. Cambridge University Press, 1989. ISBN 9780511623257. doi:[10.1017/cbo9780511623257](https://doi.org/10.1017/cbo9780511623257). URL <http://dx.doi.org/10.1017/cbo9780511623257>.
- Daniel J. Amit, Hanoch Gutfreund, and H. Sompolinsky. Spin-glass models of neural networks. *Physical Review A*, 32(2):1007–1018, 8 1985. ISSN 0556-2791. doi:[10.1103/physreva.32.1007](https://doi.org/10.1103/physreva.32.1007). URL <http://dx.doi.org/10.1103/physreva.32.1007>.
- P. W. Anderson. More Is Different: Broken symmetry and the nature of the hierarchical structure of science. *Science*, 177(4047):393–396, 8 1972. ISSN 1095-9203. doi:[10.1126/science.177.4047.393](https://doi.org/10.1126/science.177.4047.393). URL <http://dx.doi.org/10.1126/science.177.4047.393>.
- Mel Andrews. The math is not the territory: navigating the free energy principle. *Biology & Philosophy*, 36(3), 2021. ISSN 1572-8404. doi:[10.1007/s10539-021-09807-0](https://doi.org/10.1007/s10539-021-09807-0). URL <http://dx.doi.org/10.1007/s10539-021-09807-0>.
- P Ao. Potential in stochastic differential equations: novel construction. *Journal of Physics A: Mathematical and General*, 37(3):L25–L30, 1 2004. ISSN 1361-6447. doi:[10.1088/0305-4470/37/3/L01](https://doi.org/10.1088/0305-4470/37/3/L01). URL <http://dx.doi.org/10.1088/0305-4470/37/3/L01>.
- Abhay Ashtekar and Troy A. Schilling. *Geometrical Formulation of Quantum Mechanics*, pages 23–65. Springer New York, 1999. ISBN 9781461214229. doi:[10.1007/978-1-4612-1422-9_3](https://doi.org/10.1007/978-1-4612-1422-9_3). URL http://dx.doi.org/10.1007/978-1-4612-1422-9_3.
- Anastasios P. Athanasiadis, Marmaduke M. Woodman, Meysam Hashemi, Paul Triebkorn, Spase Petkoski, and Viktor K. Jirsa. A deformable attractor manifold organizes human resting-state brain dynamics. 12 2025. doi:[10.64898/2025.12.02.691788](https://doi.org/10.64898/2025.12.02.691788). URL <http://dx.doi.org/10.64898/2025.12.02.691788>.
- Hagai Attias. Planning by probabilistic inference. In *Proceedings of the Ninth International Workshop on Artificial Intelligence and Statistics (AISTATS)*, 2003.
- Bernard J. Baars. *A Cognitive Theory of Consciousness*. Cambridge University Press, 1988.

- Adriano Barra, Alberto Bernacchia, Enrica Santucci, and Pierluigi Contucci. On the equivalence of Hopfield networks and Boltzmann Machines. *Neural Networks*, 34:1–9, 10 2012. ISSN 0893-6080. doi:[10.1016/j.neunet.2012.06.003](https://doi.org/10.1016/j.neunet.2012.06.003). URL <http://dx.doi.org/10.1016/j.neunet.2012.06.003>.
- Andre M. Bastos, W. Martin Usrey, Rick A. Adams, George R. Mangun, Pascal Fries, and Karl J. Friston. Canonical Microcircuits for Predictive Coding. *Neuron*, 76(4):695–711, 11 2012. ISSN 0896-6273. doi:[10.1016/j.neuron.2012.10.038](https://doi.org/10.1016/j.neuron.2012.10.038). URL <http://dx.doi.org/10.1016/j.neuron.2012.10.038>.
- Pietro Berkes, Gergő Orbán, Máté Lengyel, and József Fiser. Spontaneous Cortical Activity Reveals Hallmarks of an Optimal Internal Model of the Environment. *Science*, 331(6013):83–87, 1 2011. ISSN 1095-9203. doi:[10.1126/science.1195870](https://doi.org/10.1126/science.1195870). URL <http://dx.doi.org/10.1126/science.1195870>.
- Michael Victor Berry. Quantal phase factors accompanying adiabatic changes. *Proceedings of the Royal Society of London. A. Mathematical and Physical Sciences*, 392(1802):45–57, 3 1984. ISSN 2053-9169. doi:[10.1098/rspa.1984.0023](https://doi.org/10.1098/rspa.1984.0023). URL <http://dx.doi.org/10.1098/rspa.1984.0023>.
- Guo-qiang Bi and Mu-ming Poo. Synaptic Modifications in Cultured Hippocampal Neurons: Dependence on Spike Timing, Synaptic Strength, and Postsynaptic Cell Type. *The Journal of Neuroscience*, 18(24):10464–10472, 12 1998. ISSN 1529-2401. doi:[10.1523/jneurosci.18-24-10464.1998](https://doi.org/10.1523/jneurosci.18-24-10464.1998). URL <http://dx.doi.org/10.1523/jneurosci.18-24-10464.1998>.
- Martin Biehl, Felix A. Pollock, and Ryota Kanai. A Technical Critique of Some Parts of the Free Energy Principle. *Entropy*, 23(3):293, 2 2021. ISSN 1099-4300. doi:[10.3390/e23030293](https://doi.org/10.3390/e23030293). URL <http://dx.doi.org/10.3390/e23030293>.
- William A. Borders, Ahmed Z. Pervaiz, Shunsuke Fukami, Kerem Y. Camsari, Hideo Ohno, and Supriyo Datta. Integer factorization using stochastic magnetic tunnel junctions. *Nature*, 573(7774):390–393, 2019. ISSN 1476-4687. doi:[10.1038/s41586-019-1557-9](https://doi.org/10.1038/s41586-019-1557-9). URL <http://dx.doi.org/10.1038/s41586-019-1557-9>.
- Matthew Botvinick and Marc Toussaint. Planning as inference. *Trends in Cognitive Sciences*, 16(10):485–488, 10 2012. ISSN 1364-6613. doi:[10.1016/j.tics.2012.08.006](https://doi.org/10.1016/j.tics.2012.08.006). URL <http://dx.doi.org/10.1016/j.tics.2012.08.006>.
- Mark E. Bouton. Context and Behavioral Processes in Extinction: Table 1. *Learning & Memory*, 11(5):485–494, 2004. ISSN 1549-5485. doi:[10.1101/lm.78804](https://doi.org/10.1101/lm.78804). URL <http://dx.doi.org/10.1101/lm.78804>.
- Michael Breakspear. Dynamic models of large-scale brain activity. *Nature Neuroscience*, 20(3):340–352, 2 2017. ISSN 1546-1726. doi:[10.1038/nn.4497](https://doi.org/10.1038/nn.4497). URL <http://dx.doi.org/10.1038/nn.4497>.
- Jelle Bruineberg, Krzysztof Dołęga, Joe Dewhurst, and Manuel Baltieri. The Emperor’s New Markov Blankets. *Behavioral and Brain Sciences*, 45, 10 2021. ISSN 1469-1825. doi:[10.1017/s0140525x21002351](https://doi.org/10.1017/s0140525x21002351). URL <http://dx.doi.org/10.1017/s0140525x21002351>.
- Ugo Bruno, Matteo Farronato, Alberto Salleo, Simone Fabiano, and Francesca Santoro. Figures of merit for neuromorphic devices through biological emulation efficacy. *Nature Reviews Electrical Engineering*, 3(9):579–595, 8 2026. ISSN 2948-1201. doi:[10.1038/s44287-026-00321-7](https://doi.org/10.1038/s44287-026-00321-7). URL <http://dx.doi.org/10.1038/s44287-026-00321-7>.
- Christian Büchel, Stephan Geuter, Christian Sprenger, and Falk Eippert. Placebo Analgesia: A Predictive Coding Perspective. *Neuron*, 81(6):1223–1239, 3 2014. ISSN 0896-6273. doi:[10.1016/j.neuron.2014.02.042](https://doi.org/10.1016/j.neuron.2014.02.042). URL <http://dx.doi.org/10.1016/j.neuron.2014.02.042>.
- Lars Buesing, Johannes Bill, Bernhard Nessler, and Wolfgang Maass. Neural Dynamics as Sampling: A Model for Stochastic Computation in Recurrent Networks of Spiking Neurons. *PLoS Computational Biology*, 7(11):e1002211, 11 2011. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1002211](https://doi.org/10.1371/journal.pcbi.1002211). URL <http://dx.doi.org/10.1371/journal.pcbi.1002211>.
- Jerome R. Busemeyer and Peter D. Bruza. *Quantum Models of Cognition and Decision*. Cambridge University Press, 2012. ISBN 9781107419889. doi:[10.1017/cbo9780511997716](https://doi.org/10.1017/cbo9780511997716). URL <http://dx.doi.org/10.1017/cbo9780511997716>.
- György Buzsáki. Theta Oscillations in the Hippocampus. *Neuron*, 33(3):325–340, 1 2002. ISSN 0896-6273. doi:[10.1016/s0896-6273\(02\)00586-x](https://doi.org/10.1016/s0896-6273(02)00586-x). URL [http://dx.doi.org/10.1016/s0896-6273\(02\)00586-x](http://dx.doi.org/10.1016/s0896-6273(02)00586-x).
- Joana Cabral, Etienne Hugues, Olaf Sporns, and Gustavo Deco. Role of local network oscillations in resting-state functional connectivity. *NeuroImage*, 57(1):130–139, 2011. ISSN 1053-8119. doi:[10.1016/j.neuroimage.2011.04.010](https://doi.org/10.1016/j.neuroimage.2011.04.010). URL <http://dx.doi.org/10.1016/j.neuroimage.2011.04.010>.

- Kerem Y. Camsari, Brian M. Sutton, and Supriyo Datta. p-bits for probabilistic spin logic. *Applied Physics Reviews*, 6(1), 3 2019. ISSN 1931-9401. doi:[10.1063/1.5055860](https://doi.org/10.1063/1.5055860). URL <http://dx.doi.org/10.1063/1.5055860>.
- Jack Carr. *Applications of Centre Manifold Theory*. Springer US, 1981. ISBN 9781461259299. doi:[10.1007/978-1-4612-5929-9](https://doi.org/10.1007/978-1-4612-5929-9). URL <http://dx.doi.org/10.1007/978-1-4612-5929-9>.
- Ariel Caticha. Entropic dynamics, time and quantum theory. *Journal of Physics A: Mathematical and Theoretical*, 44(22):225303, 5 2011. ISSN 1751-8121. doi:[10.1088/1751-8113/44/22/225303](https://doi.org/10.1088/1751-8113/44/22/225303). URL <http://dx.doi.org/10.1088/1751-8113/44/22/225303>.
- Sarthak Chandra, Sugandha Sharma, Rishidev Chaudhuri, and Ila Fiete. Episodic and associative memory from spatial scaffolds in the hippocampus. *Nature*, 638(8051):739–751, 1 2025. ISSN 1476-4687. doi:[10.1038/s41586-024-08392-y](https://doi.org/10.1038/s41586-024-08392-y). URL <http://dx.doi.org/10.1038/s41586-024-08392-y>.
- David G. Clark. Transient dynamics of associative memory models. *Physical Review E*, 113(5), 5 2026. ISSN 2470-0053. doi:[10.1103/42y2-bsh1](https://doi.org/10.1103/42y2-bsh1). URL <http://dx.doi.org/10.1103/42y2-bsh1>.
- Robert E. Clark and Larry R. Squire. Classical Conditioning and Brain Systems: The Role of Awareness. *Science*, 280(5360):77–81, 4 1998. ISSN 1095-9203. doi:[10.1126/science.280.5360.77](https://doi.org/10.1126/science.280.5360.77). URL <http://dx.doi.org/10.1126/science.280.5360.77>.
- Michael W Cole, Takuya Ito, Danielle S Bassett, and Douglas H Schultz. Activity flow over resting-state networks shapes cognitive task activations. *Nature Neuroscience*, 19(12):1718–1726, 10 2016. ISSN 1546-1726. doi:[10.1038/nn.4406](https://doi.org/10.1038/nn.4406). URL <http://dx.doi.org/10.1038/nn.4406>.
- Patrick J. Coles, Collin Szczepanski, Denis Melanson, Kaelan Donatella, Antonio J. Martinez, and Faris Sbahi. Thermodynamic AI and the Fluctuation Frontier. In *2023 IEEE International Conference on Rebooting Computing (ICRC)*, pages 1–10. IEEE, 12 2023. doi:[10.1109/icrc60800.2023.10386858](https://doi.org/10.1109/icrc60800.2023.10386858). URL <http://dx.doi.org/10.1109/icrc60800.2023.10386858>.
- Laura Lee Colgin, Edvard I. Moser, and May-Britt Moser. Understanding memory through hippocampal remapping. *Trends in Neurosciences*, 31(9):469–477, 2008. ISSN 0166-2236. doi:[10.1016/j.tins.2008.06.008](https://doi.org/10.1016/j.tins.2008.06.008). URL <http://dx.doi.org/10.1016/j.tins.2008.06.008>.
- Tom Conte, Erik DeBenedictis, Natesh Ganesh, Todd Hylton, John Paul Strachan, R. Stanley Williams, Alexander Alemi, Lee Altenberg, Gavin Crooks, James Crutchfield, Lidia del Rio, Josh Deutsch, Michael DeWeese, Khari Douglas, Massimiliano Esposito, Michael Frank, Robert Fry, Peter Harsha, Mark Hill, Christopher Kello, Jeff Krichmar, Suhas Kumar, Shih-Chii Liu, Seth Lloyd, Matteo Marsili, Ilya Nemenman, Alex Nugent, Norman Packard, Dana Randall, Peter Sadowski, Narayana Santhanam, Robert Shaw, Adam Stieg, Elan Stopnitzky, Christof Teuscher, Chris Watkins, David Wolpert, Joshua Yang, and Yan Yufik. Thermodynamic Computing, 2019. URL <https://arxiv.org/abs/1911.01968>.
- Philip R. Corlett, Guillermo Horga, Paul C. Fletcher, Ben Alderson-Day, Katharina Schmack, and Albert R. Powers. Hallucinations and Strong Priors. *Trends in Cognitive Sciences*, 23(2):114–127, 2 2019. ISSN 1364-6613. doi:[10.1016/j.tics.2018.12.001](https://doi.org/10.1016/j.tics.2018.12.001). URL <http://dx.doi.org/10.1016/j.tics.2018.12.001>.
- Gustavo Deco and Viktor K. Jirsa. Ongoing Cortical Activity at Rest: Criticality, Multistability, and Ghost Attractors. *The Journal of Neuroscience*, 32(10):3366–3375, 3 2012. ISSN 1529-2401. doi:[10.1523/jneurosci.2523-11.2012](https://doi.org/10.1523/jneurosci.2523-11.2012). URL <http://dx.doi.org/10.1523/jneurosci.2523-11.2012>.
- Gustavo Deco, Viktor K. Jirsa, Peter A. Robinson, Michael Breakspear, and Karl Friston. The Dynamic Brain: From Spiking Neurons to Neural Masses and Cortical Fields. *PLoS Computational Biology*, 4(8):e1000092, aug 29 2008. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1000092](https://doi.org/10.1371/journal.pcbi.1000092). URL <http://dx.doi.org/10.1371/journal.pcbi.1000092>.
- Gustavo Deco, Viktor K. Jirsa, and Anthony R. McIntosh. Emerging concepts for the dynamical organization of resting-state activity in the brain. *Nature Reviews Neuroscience*, 12(1):43–56, 12 2010. ISSN 1471-0048. doi:[10.1038/nrn2961](https://doi.org/10.1038/nrn2961). URL <http://dx.doi.org/10.1038/nrn2961>.
- Stanislas Dehaene, Michel Kerszberg, and Jean-Pierre Changeux. A neuronal model of a global workspace in effortful cognitive tasks. *Proceedings of the National Academy of Sciences*, 95(24):14529–14534, 11 1998. ISSN 1091-6490. doi:[10.1073/pnas.95.24.14529](https://doi.org/10.1073/pnas.95.24.14529). URL <http://dx.doi.org/10.1073/pnas.95.24.14529>.
- Adrien Doerig, Rowan P. Sommers, Katja Seeliger, Blake Richards, Jenann Ismael, Grace W. Lindsay, Konrad P. Kording, Talia Konkle, Marcel A. J. van Gerven, Nikolaus Kriegeskorte, and Tim C. Kietzmann. The neuroconnectionist research programme. *Nature Reviews Neuroscience*, 24(7):431–450, 5 2023. ISSN 1471-0048. doi:[10.1038/s41583-023-00705-w](https://doi.org/10.1038/s41583-023-00705-w). URL <http://dx.doi.org/10.1038/s41583-023-00705-w>.

- Peter Dombrowski. On the Geometry of the Tangent Bundle. *crll*, 1962(210):73–88, 1962. ISSN 0075-4102. doi:[10.1515/crll.1962.210.73](https://doi.org/10.1515/crll.1962.210.73). URL <http://dx.doi.org/10.1515/crll.1962.210.73>.
- Nico U. F. Dosenbach, Marcus E. Raichle, and Evan M. Gordon. The brain’s action-mode network. *Nature Reviews Neuroscience*, 26(3):158–168, 1 2025. ISSN 1471-0048. doi:[10.1038/s41583-024-00895-x](https://doi.org/10.1038/s41583-024-00895-x). URL <http://dx.doi.org/10.1038/s41583-024-00895-x>.
- Daniel Durstewitz, Georgia Koppe, and Max Ingo Thurm. Reconstructing computational system dynamics from neural data with recurrent neural networks. *Nature Reviews Neuroscience*, 24(11):693–710, 10 2023. ISSN 1471-0048. doi:[10.1038/s41583-023-00740-7](https://doi.org/10.1038/s41583-023-00740-7). URL <http://dx.doi.org/10.1038/s41583-023-00740-7>.
- Rodrigo Echeveste, Laurence Aitchison, Guillaume Hennequin, and Máté Lengyel. Cortical-like dynamics in recurrent circuits optimized for sampling-based probabilistic inference. *Nature Neuroscience*, 23(9):1138–1149, aug 10 2020. ISSN 1097-6256. doi:[10.1038/s41593-020-0671-1](https://doi.org/10.1038/s41593-020-0671-1). URL <http://dx.doi.org/10.1038/s41593-020-0671-1>.
- Simon B. Eickhoff, B. T. Thomas Yeo, and Sarah Genon. Imaging-based parcellations of the human brain. *Nature Reviews Neuroscience*, 19(11):672–686, 10 2018. ISSN 1471-0048. doi:[10.1038/s41583-018-0071-7](https://doi.org/10.1038/s41583-018-0071-7). URL <http://dx.doi.org/10.1038/s41583-018-0071-7>.
- Jeremy L. England. Dissipative adaptation in driven self-assembly. *Nature Nanotechnology*, 10(11):919–923, 11 2015. ISSN 1748-3395. doi:[10.1038/nnano.2015.250](https://doi.org/10.1038/nnano.2015.250). URL <http://dx.doi.org/10.1038/nnano.2015.250>.
- Robert Englert, Balint Kincses, Raviteja Kotikalapudi, Giuseppe Gallitto, Jialin Li, Kevin Hoffschlag, Choong-Wan Woo, Tor D Wager, Dagmar Timmann, Ulrike Bingel, and Tamas Spisak. Functional connectivity-based attractor dynamics of the human brain in rest, task, and disease. *eLife*, 13, mar 18 2026. ISSN 2050-084X. doi:[10.7554/elife.98725](https://doi.org/10.7554/elife.98725). URL <http://dx.doi.org/10.7554/elife.98725>.
- Takahiro Ezaki, Takamitsu Watanabe, Masayuki Ohzeki, and Naoki Masuda. Energy landscape analysis of neuroimaging data. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 375(2096):20160287, 5 2017. ISSN 1471-2962. doi:[10.1098/rsta.2016.0287](https://doi.org/10.1098/rsta.2016.0287). URL <http://dx.doi.org/10.1098/rsta.2016.0287>.
- Alberto Fachechi, Elena Agliari, and Adriano Barra. Dreaming neural networks: Forgetting spurious memories and reinforcing pure ones. *Neural Networks*, 112:24–40, 4 2019. ISSN 0893-6080. doi:[10.1016/j.neunet.2019.01.006](https://doi.org/10.1016/j.neunet.2019.01.006). URL <http://dx.doi.org/10.1016/j.neunet.2019.01.006>.
- A. Aldo Faisal, Luc P. J. Selen, and Daniel M. Wolpert. Noise in the nervous system. *Nature Reviews Neuroscience*, 9(4):292–303, 4 2008. ISSN 1471-0048. doi:[10.1038/nrn2258](https://doi.org/10.1038/nrn2258). URL <http://dx.doi.org/10.1038/nrn2258>.
- Diego Fasoli, Ludovico Coletta, Daniel Gutierrez-Barragan, Silvia Gini, Alessandro Gozzi, and Stefano Panzeri. Attractor dynamics of a whole-cortex network model predicts emergence and structure of fMRI co-activation patterns in the mouse brain. *PLOS Computational Biology*, 22(2):e1013995, 2 2026. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1013995](https://doi.org/10.1371/journal.pcbi.1013995). URL <http://dx.doi.org/10.1371/journal.pcbi.1013995>.
- Harriet Feldman and Karl J. Friston. Attention, Uncertainty, and Free-Energy. *Frontiers in Human Neuroscience*, 4, 2010. ISSN 1662-5161. doi:[10.3389/fnhum.2010.00215](https://doi.org/10.3389/fnhum.2010.00215). URL <http://dx.doi.org/10.3389/fnhum.2010.00215>.
- Chris Fields and Michael Levin. Scalefree Biology: Integrating Evolutionary and Developmental Thinking. *BioEssays*, 42(8), 2020. ISSN 1521-1878. doi:[10.1002/bies.201900228](https://doi.org/10.1002/bies.201900228). URL <http://dx.doi.org/10.1002/bies.201900228>.
- Chris Fields and Michael Levin. Competency in Navigating Arbitrary Spaces as an Invariant for Analyzing Cognition in Diverse Embodiments. *Entropy*, 24(6):819, 2022. ISSN 1099-4300. doi:[10.3390/e24060819](https://doi.org/10.3390/e24060819). URL <http://dx.doi.org/10.3390/e24060819>.
- Chris Fields, James F Glazebrook, and Michael Levin. Minimal physicalism as a scale-free substrate for cognition and consciousness. *Neuroscience of Consciousness*, 2021(2), 5 2021. ISSN 2057-2107. doi:[10.1093/nc/niab013](https://doi.org/10.1093/nc/niab013). URL <http://dx.doi.org/10.1093/nc/niab013>.
- Arseny Finkelstein, Lorenzo Fontolan, Michael N. Economo, Nuo Li, Sandro Romani, and Karel Svoboda. Attractor dynamics gate cortical information flow during decision-making. *Nature Neuroscience*, 24(6):843–850, 4 2021. ISSN 1546-1726. doi:[10.1038/s41593-021-00840-6](https://doi.org/10.1038/s41593-021-00840-6). URL <http://dx.doi.org/10.1038/s41593-021-00840-6>.

- József Fiser, Pietro Berkes, Gergő Orbán, and Máté Lengyel. Statistically optimal perception and learning: from behavior to neural representations. *Trends in Cognitive Sciences*, 14(3):119–130, 3 2010. ISSN 1364-6613. doi:[10.1016/j.tics.2010.01.003](https://doi.org/10.1016/j.tics.2010.01.003). URL <http://dx.doi.org/10.1016/j.tics.2010.01.003>.
- Timo Flesch, Keno Juechems, Tsvetomira Dumbalska, Andrew Saxe, and Christopher Summerfield. Orthogonal representations for robust context-dependent task performance in brains and neural networks. *Neuron*, 110(24):4212–4219, 12 2022. ISSN 0896-6273. doi:[10.1016/j.neuron.2022.12.004](https://doi.org/10.1016/j.neuron.2022.12.004). URL <http://dx.doi.org/10.1016/j.neuron.2022.12.004>.
- Herta Flor, Lone Nikolajsen, and Troels Staehelin Jensen. Phantom limb pain: a case of maladaptive CNS plasticity? *Nature Reviews Neuroscience*, 7(11):873–881, 11 2006. ISSN 1471-0048. doi:[10.1038/nrn1991](https://doi.org/10.1038/nrn1991). URL <http://dx.doi.org/10.1038/nrn1991>.
- P. Földiák. Forming sparse representations by local anti-Hebbian learning. *Biological Cybernetics*, 64(2):165–170, 12 1990. ISSN 0340-1200. doi:[10.1007/bf02331346](https://doi.org/10.1007/bf02331346). URL <http://dx.doi.org/10.1007/bf02331346>.
- Michael D. Fox, Abraham Z. Snyder, Justin L. Vincent, Maurizio Corbetta, David C. Van Essen, and Marcus E. Raichle. The human brain is intrinsically organized into dynamic, anticorrelated functional networks. *Proceedings of the National Academy of Sciences*, 102(27):9673–9678, 2005. ISSN 1091-6490. doi:[10.1073/pnas.0504136102](https://doi.org/10.1073/pnas.0504136102). URL <http://dx.doi.org/10.1073/pnas.0504136102>.
- R. French. Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3(4):128–135, 4 1999. ISSN 1364-6613. doi:[10.1016/s1364-6613\(99\)01294-2](https://doi.org/10.1016/s1364-6613(99)01294-2). URL [http://dx.doi.org/10.1016/s1364-6613\(99\)01294-2](http://dx.doi.org/10.1016/s1364-6613(99)01294-2).
- Karl Friston. A theory of cortical responses. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 360(1456):815–836, apr 29 2005. ISSN 0962-8436. doi:[10.1098/rstb.2005.1622](https://doi.org/10.1098/rstb.2005.1622). URL <http://dx.doi.org/10.1098/rstb.2005.1622>.
- Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 1 2010. ISSN 1471-0048. doi:[10.1038/nrn2787](https://doi.org/10.1038/nrn2787). URL <http://dx.doi.org/10.1038/nrn2787>.
- Karl Friston. A free energy principle for a particular physics, 2019. URL <https://arxiv.org/abs/1906.10184>.
- Karl Friston and Ping Ao. Free Energy, Value, and Attractors. *Computational and Mathematical Methods in Medicine*, 2012:1–27, 2012. ISSN 1748-6718. doi:[10.1155/2012/937860](https://doi.org/10.1155/2012/937860). URL <http://dx.doi.org/10.1155/2012/937860>.
- Karl Friston, Francesco Rigoli, Dimitri Ognibene, Christoph Mathys, Thomas Fitzgerald, and Giovanni Pezzulo. Active inference and epistemic value. *Cognitive Neuroscience*, 6(4):187–214, 3 2015. ISSN 1758-8936. doi:[10.1080/17588928.2015.1020053](https://doi.org/10.1080/17588928.2015.1020053). URL <http://dx.doi.org/10.1080/17588928.2015.1020053>.
- Karl Friston, Lancelot Da Costa, Noor Sajid, Conor Heins, Kai Ueltzhöffer, Grigorios A. Pavliotis, and Thomas Parr. The free energy principle made simpler but not too simple. *Physics Reports*, 1024:1–29, 6 2023a. ISSN 0370-1573. doi:[10.1016/j.physrep.2023.07.001](https://doi.org/10.1016/j.physrep.2023.07.001). URL <http://dx.doi.org/10.1016/j.physrep.2023.07.001>.
- Karl Friston, Lancelot Da Costa, Dalton A.R. Sakthivadivel, Conor Heins, Grigorios A. Pavliotis, Maxwell Ramstead, and Thomas Parr. Path integrals, particular kinds, and strange things. *Physics of Life Reviews*, 47:35–62, 12 2023b. ISSN 1571-0645. doi:[10.1016/j.plrev.2023.08.016](https://doi.org/10.1016/j.plrev.2023.08.016). URL <http://dx.doi.org/10.1016/j.plrev.2023.08.016>.
- Karl Friston, Conor Heins, Tim Verbelen, Lancelot Da Costa, Tommaso Salvatori, Dimitrije Markovic, Alexander Tschantz, Magnus Koudahl, Christopher Buckley, and Thomas Parr. From pixels to planning: scale-free active inference. *Frontiers in Network Physiology*, 5, 2025. ISSN 2674-0109. doi:[10.3389/fnetp.2025.1521963](https://doi.org/10.3389/fnetp.2025.1521963). URL <http://dx.doi.org/10.3389/fnetp.2025.1521963>.
- Karl J. Friston, Jean Daunizeau, and Stefan J. Kiebel. Reinforcement Learning or Active Inference? *PLoS ONE*, 4(7):e6421, jul 29 2009. ISSN 1932-6203. doi:[10.1371/journal.pone.0006421](https://doi.org/10.1371/journal.pone.0006421). URL <http://dx.doi.org/10.1371/journal.pone.0006421>.
- Karl J. Friston, Joshua Kahan, Adeel Razi, Klaas Enno Stephan, and Olaf Sporns. On nodes and modes in resting state fMRI. *NeuroImage*, 99:533–547, 10 2014. ISSN 1053-8119. doi:[10.1016/j.neuroimage.2014.05.056](https://doi.org/10.1016/j.neuroimage.2014.05.056). URL <http://dx.doi.org/10.1016/j.neuroimage.2014.05.056>.
- Karl J. Friston, Marco Lin, Christopher D. Frith, Giovanni Pezzulo, J. Allan Hobson, and Sasha Ondobaka. Active Inference, Curiosity and Insight. *Neural Computation*, 29(10):2633–2683, 10 2017. ISSN 1530-888X. doi:[10.1162/neco_a_00999](https://doi.org/10.1162/neco_a_00999). URL http://dx.doi.org/10.1162/neco_a_00999.

- Karl J. Friston, Erik D. Fagerholm, Tahereh S. Zarghami, Thomas Parr, Inês Hipólito, Loïc Magrou, and Adeel Razi. Parcels and particles: Markov blankets in the brain. *Network Neuroscience*, 5(1):211–251, 1 2021. ISSN 2472-1751. doi:[10.1162/netn_a_00175](https://doi.org/10.1162/netn_a_00175). URL http://dx.doi.org/10.1162/netn_a_00175.
- K.J. Friston, L. Harrison, and W. Penny. Dynamic causal modelling. *NeuroImage*, 19(4):1273–1302, 8 2003. ISSN 1053-8119. doi:[10.1016/s1053-8119\(03\)00202-7](https://doi.org/10.1016/s1053-8119(03)00202-7). URL [http://dx.doi.org/10.1016/s1053-8119\(03\)00202-7](http://dx.doi.org/10.1016/s1053-8119(03)00202-7).
- Samuel J. Gershman, David M. Blei, and Yael Niv. Context, learning, and extinction. *Psychological Review*, 117(1):197–209, 2010. ISSN 0033-295X. doi:[10.1037/a0017808](https://doi.org/10.1037/a0017808). URL <http://dx.doi.org/10.1037/a0017808>.
- Samuel J Gershman, Marie-H Monfils, Kenneth A Norman, and Yael Niv. The computational nature of memory modification. *eLife*, 6, 3 2017. ISSN 2050-084X. doi:[10.7554/elife.23763](https://doi.org/10.7554/elife.23763). URL <http://dx.doi.org/10.7554/elife.23763>.
- Wulfram Gerstner, Marco Lehmann, Vasiliki Liakoni, Dane Corneil, and Johanni Brea. Eligibility Traces and Plasticity on Behavioral Time Scales: Experimental Support of NeoHebbian Three-Factor Learning Rules. *Frontiers in Neural Circuits*, 12, 2018. ISSN 1662-5110. doi:[10.3389/fncir.2018.00053](https://doi.org/10.3389/fncir.2018.00053). URL <http://dx.doi.org/10.3389/fncir.2018.00053>.
- Matthieu Gilson, Ruben Moreno-Bote, Adrián Ponce-Alvarez, Petra Ritter, and Gustavo Deco. Estimation of Directed Effective Connectivity from fMRI Functional Connectivity Hints at Asymmetries of Cortical Connectome. *PLOS Computational Biology*, 12(3):e1004762, 3 2016. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1004762](https://doi.org/10.1371/journal.pcbi.1004762). URL <http://dx.doi.org/10.1371/journal.pcbi.1004762>.
- P.S Goldman-Rakic. Cellular basis of working memory. *Neuron*, 14(3):477–485, 3 1995. ISSN 0896-6273. doi:[10.1016/0896-6273\(95\)90304-6](https://doi.org/10.1016/0896-6273(95)90304-6). URL [http://dx.doi.org/10.1016/0896-6273\(95\)90304-6](http://dx.doi.org/10.1016/0896-6273(95)90304-6).
- Sebastian Goldt and Udo Seifert. Stochastic Thermodynamics of Learning. *Physical Review Letters*, 118(1), 1 2017. ISSN 1079-7114. doi:[10.1103/physrevlett.118.010601](https://doi.org/10.1103/physrevlett.118.010601). URL <http://dx.doi.org/10.1103/physrevlett.118.010601>.
- Mathieu Golos, Viktor Jirsa, and Emmanuel Daucé. Multistability in Large Scale Models of Brain Activity. *PLOS Computational Biology*, 11(12):e1004644, 12 2015. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1004644](https://doi.org/10.1371/journal.pcbi.1004644). URL <http://dx.doi.org/10.1371/journal.pcbi.1004644>.
- Evan M. Gordon, Timothy O. Laumann, Babatunde Adeyemo, Jeremy F. Huckins, William M. Kelley, and Steven E. Petersen. Generation and Evaluation of a Cortical Area Parcellation from Resting-State Correlations. *Cerebral Cortex*, 26(1):288–303, oct 14 2014. ISSN 1047-3211. doi:[10.1093/cercor/bhu239](https://doi.org/10.1093/cercor/bhu239). URL <http://dx.doi.org/10.1093/cercor/bhu239>.
- Evan M. Gordon, Timothy O. Laumann, Adrian W. Gilmore, Dillan J. Newbold, Deanna J. Greene, Jeffrey J. Berg, Mario Ortega, Catherine Hoyt-Drazen, Caterina Gratton, Haixin Sun, Jacqueline M. Hampton, Rebecca S. Coalson, Annie L. Nguyen, Kathleen B. McDermott, Joshua S. Shimony, Abraham Z. Snyder, Bradley L. Schlaggar, Steven E. Petersen, Steven M. Nelson, and Nico U.F. Dosenbach. Precision Functional Mapping of Individual Human Brains. *Neuron*, 95(4):791–807.e7, 8 2017. ISSN 0896-6273. doi:[10.1016/j.neuron.2017.07.011](https://doi.org/10.1016/j.neuron.2017.07.011). URL <http://dx.doi.org/10.1016/j.neuron.2017.07.011>.
- Giorgio Gosti, Edoardo Milanetti, Viola Folli, Francesco de Pasquale, Marco Leonetti, Maurizio Corbetta, Giancarlo Ruocco, and Stefania Della Penna. A recurrent Hopfield network for estimating meso-scale effective connectivity in MEG. *Neural Networks*, 170:72–93, 2 2024. ISSN 0893-6080. doi:[10.1016/j.neunet.2023.11.027](https://doi.org/10.1016/j.neunet.2023.11.027). URL <http://dx.doi.org/10.1016/j.neunet.2023.11.027>.
- Jacqueline Gottlieb and Pierre-Yves Oudeyer. Towards a neuroscience of active sampling and curiosity. *Nature Reviews Neuroscience*, 19(12):758–770, 11 2018. ISSN 1471-0048. doi:[10.1038/s41583-018-0078-0](https://doi.org/10.1038/s41583-018-0078-0). URL <http://dx.doi.org/10.1038/s41583-018-0078-0>.
- Miroslav Grmela and Hans Christian Öttinger. Dynamics and thermodynamics of complex fluids. I. Development of a general formalism. *Physical Review E*, 56(6):6620–6632, 12 1997. ISSN 1095-3787. doi:[10.1103/physreve.56.6620](https://doi.org/10.1103/physreve.56.6620). URL <http://dx.doi.org/10.1103/physreve.56.6620>.
- Shi Gu, Matthew Cieslak, Benjamin Baird, Sarah F. Muldoon, Scott T. Grafton, Fabio Pasqualetti, and Danielle S. Bassett. The Energy Landscape of Neurophysiological Activity Implicit in Brain Network Structure. *Scientific Reports*, 8(1), 2 2018. ISSN 2045-2322. doi:[10.1038/s41598-018-20123-8](https://doi.org/10.1038/s41598-018-20123-8). URL <http://dx.doi.org/10.1038/s41598-018-20123-8>.
- Raia Hadsell, Dushyant Rao, Andrei A. Rusu, and Razvan Pascanu. Embracing Change: Continual Learning in Deep Neural Networks. *Trends in Cognitive Sciences*, 24(12):1028–1040, 12 2020. ISSN 1364-6613. doi:[10.1016/j.tics.2020.09.004](https://doi.org/10.1016/j.tics.2020.09.004). URL <http://dx.doi.org/10.1016/j.tics.2020.09.004>.

- Hermann Haken. *Synergetics: An Introduction*. Springer Berlin Heidelberg, 1983. ISBN 9783642883385. doi:[10.1007/978-3-642-88338-5](https://doi.org/10.1007/978-3-642-88338-5). URL <http://dx.doi.org/10.1007/978-3-642-88338-5>.
- Fran Hancock, Fernando E. Rosas, Andrea I. Luppi, Mengsen Zhang, Pedro A. M. Mediano, Joana Cabral, Gustavo Deco, Morten L. Kringelbach, Michael Breakspear, J. A. Scott Kelso, and Federico E. Turkheimer. Metastability demystified — the foundational past, the pragmatic present and the promising future. *Nature Reviews Neuroscience*, 26(2):82–100, 12 2024. ISSN 1471-0048. doi:[10.1038/s41583-024-00883-1](https://doi.org/10.1038/s41583-024-00883-1). URL <http://dx.doi.org/10.1038/s41583-024-00883-1>.
- Kamel-Eddine Harabi, Tifenn Hirtzlin, Clément Turck, Elisa Vianello, Raphaël Laurent, Jacques Droulez, Pierre Bessière, Jean-Michel Portal, Marc Bocquet, and Damien Querlioz. A memristor-based Bayesian machine. *Nature Electronics*, 12 2022. ISSN 2520-1131. doi:[10.1038/s41928-022-00886-9](https://doi.org/10.1038/s41928-022-00886-9). URL <http://dx.doi.org/10.1038/s41928-022-00886-9>.
- Demis Hassabis, Dharshan Kumaran, Christopher Summerfield, and Matthew Botvinick. Neuroscience-Inspired Artificial Intelligence. *Neuron*, 95(2):245–258, 2017. ISSN 0896-6273. doi:[10.1016/j.neuron.2017.06.011](https://doi.org/10.1016/j.neuron.2017.06.011). URL <http://dx.doi.org/10.1016/j.neuron.2017.06.011>.
- Michael E. Hasselmo, Clara Bodelón, and Bradley P. Wyble. A Proposed Function for Hippocampal Theta Rhythm: Separate Phases of Encoding and Retrieval Enhance Reversal of Prior Learning. *Neural Computation*, 14(4):793–817, 4 2002. ISSN 1530-888X. doi:[10.1162/089976602317318965](https://doi.org/10.1162/089976602317318965). URL <http://dx.doi.org/10.1162/089976602317318965>.
- Stewart Heitmann and Michael Breakspear. Putting the “dynamic” back into dynamic functional connectivity. *Network Neuroscience*, 2(2):150–174, 6 2018. ISSN 2472-1751. doi:[10.1162/netn_a_00041](https://doi.org/10.1162/netn_a_00041). URL http://dx.doi.org/10.1162/netn_a_00041.
- Olivier J. Hénaff, Zoe M. Boundy-Singer, Kristof Meding, Corey M. Ziemba, and Robbe L. T. Goris. Representation of visual uncertainty through neural gain variability. *Nature Communications*, 11(1), 5 2020. ISSN 2041-1723. doi:[10.1038/s41467-020-15533-0](https://doi.org/10.1038/s41467-020-15533-0). URL <http://dx.doi.org/10.1038/s41467-020-15533-0>.
- Geoffrey E. Hinton. Training Products of Experts by Minimizing Contrastive Divergence. *Neural Computation*, 14(8):1771–1800, 8 2002. ISSN 1530-888X. doi:[10.1162/089976602760128018](https://doi.org/10.1162/089976602760128018). URL <http://dx.doi.org/10.1162/089976602760128018>.
- J.A. Hobson and K.J. Friston. Waking and dreaming consciousness: Neurobiological and functional considerations. *Progress in Neurobiology*, 98(1):82–98, 2012. ISSN 0301-0082. doi:[10.1016/j.pneurobio.2012.05.003](https://doi.org/10.1016/j.pneurobio.2012.05.003). URL <http://dx.doi.org/10.1016/j.pneurobio.2012.05.003>.
- Erik P. Hoel, Larissa Albantakis, and Giulio Tononi. Quantifying causal emergence shows that macro can beat micro. *Proceedings of the National Academy of Sciences*, 110(49):19790–19795, 11 2013. ISSN 1091-6490. doi:[10.1073/pnas.1314922110](https://doi.org/10.1073/pnas.1314922110). URL <http://dx.doi.org/10.1073/pnas.1314922110>.
- C. J. Honey, O. Sporns, L. Cammoun, X. Gigandet, J. P. Thiran, R. Meuli, and P. Hagmann. Predicting human resting-state functional connectivity from structural connectivity. *Proceedings of the National Academy of Sciences*, 106(6):2035–2040, feb 10 2009. ISSN 0027-8424. doi:[10.1073/pnas.0811168106](https://doi.org/10.1073/pnas.0811168106). URL <http://dx.doi.org/10.1073/pnas.0811168106>.
- Sara Hooker. The hardware lottery. *Communications of the ACM*, 64(12):58–65, 11 2021. ISSN 1557-7317. doi:[10.1145/3467017](https://doi.org/10.1145/3467017). URL <http://dx.doi.org/10.1145/3467017>.
- Benjamin Hoover, Zhaoyang Shi, Krishnakumar Balasubramanian, Dmitry Krotov, and Parikshit Ram. Dense Associative Memory with Epanechnikov Energy. In *Advances in Neural Information Processing Systems 38*, NeurIPS 2025, pages 131166–131193. Neural Information Processing Systems Foundation, Inc. (NeurIPS), 2025. doi:[10.52202/085713-3947](https://doi.org/10.52202/085713-3947). URL <http://dx.doi.org/10.52202/085713-3947>.
- J J Hopfield. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8):2554–2558, 4 1982. ISSN 1091-6490. doi:[10.1073/pnas.79.8.2554](https://doi.org/10.1073/pnas.79.8.2554). URL <http://dx.doi.org/10.1073/pnas.79.8.2554>.
- J. J. Hopfield, D. I. Feinstein, and R. G. Palmer. ‘Unlearning’ has a stabilizing effect in collective memories. *Nature*, 304(5922):158–159, 1983. ISSN 1476-4687. doi:[10.1038/304158a0](https://doi.org/10.1038/304158a0). URL <http://dx.doi.org/10.1038/304158a0>.
- Patrik O. Hoyer and Aapo Hyvärinen. Interpreting neural response variability as Monte Carlo sampling of the posterior. In *Advances in Neural Information Processing Systems*, volume 15, 2003.

- Hidehiko K. Inagaki, Lorenzo Fontolan, Sandro Romani, and Karel Svoboda. Discrete attractor dynamics underlies persistent activity in the frontal cortex. *Nature*, 566(7743):212–217, 2 2019. ISSN 1476-4687. doi:[10.1038/s41586-019-0919-7](https://doi.org/10.1038/s41586-019-0919-7). URL <http://dx.doi.org/10.1038/s41586-019-0919-7>.
- Amirhossein Jafarian, Peter Zeidman, Rob. C Wykes, Matthew Walker, and Karl J. Friston. Adiabatic dynamic causal modelling. *NeuroImage*, 238:118243, 9 2021. ISSN 1053-8119. doi:[10.1016/j.neuroimage.2021.118243](https://doi.org/10.1016/j.neuroimage.2021.118243). URL <http://dx.doi.org/10.1016/j.neuroimage.2021.118243>.
- E. T. Jaynes. Information Theory and Statistical Mechanics. *Physical Review*, 106(4):620–630, 5 1957. ISSN 0031-899X. doi:[10.1103/physrev.106.620](https://doi.org/10.1103/physrev.106.620). URL <http://dx.doi.org/10.1103/physrev.106.620>.
- Richard Jordan, David Kinderlehrer, and Felix Otto. The Variational Formulation of the Fokker–Planck Equation. *SIAM Journal on Mathematical Analysis*, 29(1):1–17, 1 1998. ISSN 1095-7154. doi:[10.1137/s0036141096303359](https://doi.org/10.1137/s0036141096303359). URL <http://dx.doi.org/10.1137/s0036141096303359>.
- I. Kanter and H. Sompolinsky. Associative recall of memory without errors. *Physical Review A*, 35(1):380–392, 1 1987. ISSN 0556-2791. doi:[10.1103/physreva.35.380](https://doi.org/10.1103/physreva.35.380). URL <http://dx.doi.org/10.1103/physreva.35.380>.
- David Kappel, Stefan Habenschuss, Robert Legenstein, and Wolfgang Maass. Network Plasticity as Bayesian Inference. *PLOS Computational Biology*, 11(11):e1004485, 11 2015. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1004485](https://doi.org/10.1371/journal.pcbi.1004485). URL <http://dx.doi.org/10.1371/journal.pcbi.1004485>.
- J. A. Scott Kelso. *Dynamic Patterns: The Self-Organization of Brain and Behavior*. MIT Press, 1995.
- Mikhail Khona and Ila R. Fiete. Attractor and integrator networks in the brain. *Nature Reviews Neuroscience*, 23(12):744–766, 11 2022. ISSN 1471-0048. doi:[10.1038/s41583-022-00642-0](https://doi.org/10.1038/s41583-022-00642-0). URL <http://dx.doi.org/10.1038/s41583-022-00642-0>.
- T. W. B. Kibble. Geometrization of quantum mechanics. *Communications in Mathematical Physics*, 65(2):189–201, 1979. ISSN 1432-0916. doi:[10.1007/bf01225149](https://doi.org/10.1007/bf01225149). URL <http://dx.doi.org/10.1007/bf01225149>.
- Stefan J. Kiebel, Jean Daunizeau, and Karl J. Friston. A Hierarchy of Time-Scales and the Brain. *PLoS Computational Biology*, 4(11):e1000209, 11 2008. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1000209](https://doi.org/10.1371/journal.pcbi.1000209). URL <http://dx.doi.org/10.1371/journal.pcbi.1000209>.
- Jaegwon Kim. *Mind in a Physical World: An Essay on the Mind-Body Problem and Mental Causation*. The MIT Press, 10 1998. ISBN 9780262277105. doi:[10.7551/mitpress/4629.001.0001](https://doi.org/10.7551/mitpress/4629.001.0001). URL <http://dx.doi.org/10.7551/mitpress/4629.001.0001>.
- Sung Soo Kim, Hervé Rouault, Shaul Druckmann, and Vivek Jayaraman. Ring attractor dynamics in the *Drosophila* central brain. *Science*, 356(6340):849–853, 5 2017. ISSN 1095-9203. doi:[10.1126/science.aal4835](https://doi.org/10.1126/science.aal4835). URL <http://dx.doi.org/10.1126/science.aal4835>.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 3 2017. ISSN 1091-6490. doi:[10.1073/pnas.1611835114](https://doi.org/10.1073/pnas.1611835114). URL <http://dx.doi.org/10.1073/pnas.1611835114>.
- D Kleinfeld. Sequential state generation by model neural networks. *Proceedings of the National Academy of Sciences*, 83(24):9469–9473, 12 1986. ISSN 1091-6490. doi:[10.1073/pnas.83.24.9469](https://doi.org/10.1073/pnas.83.24.9469). URL <http://dx.doi.org/10.1073/pnas.83.24.9469>.
- David C. Knull and Alexandre Pouget. The Bayesian brain: the role of uncertainty in neural coding and computation. *Trends in Neurosciences*, 27(12):712–719, 12 2004. ISSN 0166-2236. doi:[10.1016/j.tins.2004.10.007](https://doi.org/10.1016/j.tins.2004.10.007). URL <http://dx.doi.org/10.1016/j.tins.2004.10.007>.
- Georgia Koppe, Hazem Toutounji, Peter Kirsch, Stefanie Lis, and Daniel Durstewitz. Identifying nonlinear dynamical systems via generative recurrent neural networks with applications to fMRI. *PLOS Computational Biology*, 15(8):e1007263, 8 2019. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1007263](https://doi.org/10.1371/journal.pcbi.1007263). URL <http://dx.doi.org/10.1371/journal.pcbi.1007263>.
- Morten L. Kringelbach, Yonatan Sanz Perl, and Gustavo Deco. The Thermodynamics of Mind. *Trends in Cognitive Sciences*, 28(6):568–581, 2024. ISSN 1364-6613. doi:[10.1016/j.tics.2024.03.009](https://doi.org/10.1016/j.tics.2024.03.009). URL <http://dx.doi.org/10.1016/j.tics.2024.03.009>.
- Dmitry Krotov and John Hopfield. Large Associative Memory Problem in Neurobiology and Machine Learning, 2020. URL <https://arxiv.org/abs/2008.06996>.

- Dmitry Krotov and John J Hopfield. Dense Associative Memory for Pattern Recognition. 2016. doi:[10.48550/ARXIV.1606.01164](https://arxiv.org/abs/1606.01164). URL <https://arxiv.org/abs/1606.01164>.
- R Kubo. The fluctuation-dissipation theorem. *Reports on Progress in Physics*, 29(1):255–284, 1 1966. ISSN 0034-4885. doi:[10.1088/0034-4885/29/1/306](http://dx.doi.org/10.1088/0034-4885/29/1/306). URL <http://dx.doi.org/10.1088/0034-4885/29/1/306>.
- Yoshiki Kuramoto and Hiroya Nakao. On the concept of dynamical reduction: the case of coupled oscillators. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 377(2160):20190041, oct 28 2019. ISSN 1364-503X. doi:[10.1098/rsta.2019.0041](http://dx.doi.org/10.1098/rsta.2019.0041). URL <http://dx.doi.org/10.1098/rsta.2019.0041>.
- Chulan Kwon, Ping Ao, and David J. Thouless. Structure of stochastic dynamics near fixed points. *Proceedings of the National Academy of Sciences*, 102(37):13029–13033, 2005. ISSN 1091-6490. doi:[10.1073/pnas.0506347102](http://dx.doi.org/10.1073/pnas.0506347102). URL <http://dx.doi.org/10.1073/pnas.0506347102>.
- Imre Lakatos. The Methodology of Scientific Research Programmes, may 18 1978. URL <http://dx.doi.org/10.1017/cbo9780511621123>.
- Brenden M. Lake, Tomer D. Ullman, Joshua B. Tenenbaum, and Samuel J. Gershman. Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40, 11 2016. ISSN 1469-1825. doi:[10.1017/s0140525x16001837](http://dx.doi.org/10.1017/s0140525x16001837). URL <http://dx.doi.org/10.1017/s0140525x16001837>.
- Michael Levin. The Computational Boundary of a “Self”: Developmental Bioelectricity Drives Multicellularity and Scale-Free Cognition. *Frontiers in Psychology*, 10, dec 13 2019. ISSN 1664-1078. doi:[10.3389/fpsyg.2019.02688](http://dx.doi.org/10.3389/fpsyg.2019.02688). URL <http://dx.doi.org/10.3389/fpsyg.2019.02688>.
- Michael Levin. Darwin’s agential materials: evolutionary implications of multiscale competency in developmental biology. *Cellular and Molecular Life Sciences*, 80(6), 5 2023. ISSN 1420-9071. doi:[10.1007/s00018-023-04790-z](http://dx.doi.org/10.1007/s00018-023-04790-z). URL <http://dx.doi.org/10.1007/s00018-023-04790-z>.
- Alexandra Libby and Timothy J. Buschman. Rotational dynamics reduce interference between sensory and memory representations. *Nature Neuroscience*, 24(5):715–726, 4 2021. ISSN 1546-1726. doi:[10.1038/s41593-021-00821-9](http://dx.doi.org/10.1038/s41593-021-00821-9). URL <http://dx.doi.org/10.1038/s41593-021-00821-9>.
- Ekaterina Likhtik, Joseph M Stujenske, Mihir A Topiwala, Alexander Z Harris, and Joshua A Gordon. Prefrontal entrainment of amygdala activity signals safety in learned fear and innate anxiety. *Nature Neuroscience*, 17(1):106–113, nov 17 2013. ISSN 1097-6256. doi:[10.1038/nn.3582](http://dx.doi.org/10.1038/nn.3582). URL <http://dx.doi.org/10.1038/nn.3582>.
- Timothy P. Lillicrap, Adam Santoro, Luke Marris, Colin J. Akerman, and Geoffrey Hinton. Backpropagation and the brain. *Nature Reviews Neuroscience*, 21(6):335–346, 4 2020. ISSN 1471-0048. doi:[10.1038/s41583-020-0277-3](http://dx.doi.org/10.1038/s41583-020-0277-3). URL <http://dx.doi.org/10.1038/s41583-020-0277-3>.
- Yunzhe Liu, Raymond J. Dolan, Zeb Kurth-Nelson, and Timothy E.J. Behrens. Human Replay Spontaneously Reorganizes Experience. *Cell*, 178(3):640–652.e14, 2019. ISSN 0092-8674. doi:[10.1016/j.cell.2019.06.012](http://dx.doi.org/10.1016/j.cell.2019.06.012). URL <http://dx.doi.org/10.1016/j.cell.2019.06.012>.
- Christopher W. Lynn, Eli J. Cornblath, Lia Papadopoulos, Maxwell A. Bertolero, and Danielle S. Bassett. Broken detailed balance and entropy production in the human brain. *Proceedings of the National Academy of Sciences*, 118(47), 11 2021. ISSN 1091-6490. doi:[10.1073/pnas.2109889118](http://dx.doi.org/10.1073/pnas.2109889118). URL <http://dx.doi.org/10.1073/pnas.2109889118>.
- Pamela Lyon. The biogenic approach to cognition. *Cognitive Processing*, 7(1):11–29, 10 2005. ISSN 1612-4790. doi:[10.1007/s10339-005-0016-8](http://dx.doi.org/10.1007/s10339-005-0016-8). URL <http://dx.doi.org/10.1007/s10339-005-0016-8>.
- Wei Ji Ma, Jeffrey M Beck, Peter E Latham, and Alexandre Pouget. Bayesian inference with probabilistic population codes. *Nature Neuroscience*, 9(11):1432–1438, 10 2006. ISSN 1546-1726. doi:[10.1038/nn1790](http://dx.doi.org/10.1038/nn1790). URL <http://dx.doi.org/10.1038/nn1790>.
- Yi-An Ma, Tianqi Chen, and Emily B. Fox. A Complete Recipe for Stochastic Gradient MCMC, 2015. URL <https://arxiv.org/abs/1506.04696>.
- Benjamin B. Machta, Ricky Chachra, Mark K. Transtrum, and James P. Sethna. Parameter Space Compression Underlies Emergent Theories and Predictive Models. *Science*, 342(6158):604–607, 11 2013. ISSN 1095-9203. doi:[10.1126/science.1238723](http://dx.doi.org/10.1126/science.1238723). URL <http://dx.doi.org/10.1126/science.1238723>.
- Qiuyang Mang, Kaiyuan Liu, Bo Peng, Shreyas Pimpalgaonkar, Luke Zettlemoyer, Alex Dimakis, and Alvin Cheung. Humans Still Beat AI in the Long Horizon: Revisiting Test-Time Scaling in the Agent Era, 2026. URL <https://joyemang33.github.io/blog/2026/humans-dont-just-sample/>. Accessed 6 October 2026.

- Daniel S. Margulies, Satrajit S. Ghosh, Alexandros Goulas, Marcel Falkiewicz, Julia M. Huntenburg, Georg Langs, Gleb Bezgin, Simon B. Eickhoff, F. Xavier Castellanos, Michael Petrides, Elizabeth Jefferies, and Jonathan Smallwood. Situating the default-mode network along a principal gradient of macroscale cortical organization. *Proceedings of the National Academy of Sciences*, 113(44):12574–12579, 10 2016. ISSN 1091-6490. doi:[10.1073/pnas.1608282113](https://doi.org/10.1073/pnas.1608282113). URL <http://dx.doi.org/10.1073/pnas.1608282113>.
- George A. Mashour, Pieter Roelfsema, Jean-Pierre Changeux, and Stanislas Dehaene. Conscious Processing and the Global Neuronal Workspace Hypothesis. *Neuron*, 105(5):776–798, 3 2020. ISSN 0896-6273. doi:[10.1016/j.neuron.2020.01.026](https://doi.org/10.1016/j.neuron.2020.01.026). URL <http://dx.doi.org/10.1016/j.neuron.2020.01.026>.
- Christoph Mathys. A Bayesian foundation for individual learning under uncertainty. *Frontiers in Human Neuroscience*, 5, 2011. ISSN 1662-5161. doi:[10.3389/fnhum.2011.00039](https://doi.org/10.3389/fnhum.2011.00039). URL <http://dx.doi.org/10.3389/fnhum.2011.00039>.
- L. Mazzucato, G. La Camera, and A. Fontanini. Expectation-induced modulation of metastable activity underlies faster coding of sensory stimuli. *Nature Neuroscience*, 22(5):787–796, 4 2019. ISSN 1546-1726. doi:[10.1038/s41593-019-0364-9](https://doi.org/10.1038/s41593-019-0364-9). URL <http://dx.doi.org/10.1038/s41593-019-0364-9>.
- James L. McClelland, Bruce L. McNaughton, and Randall C. O'Reilly. Why there are complementary learning systems in the hippocampus and neocortex: Insights from the successes and failures of connectionist models of learning and memory. *Psychological Review*, 102(3):419–457, 1995. ISSN 0033-295X. doi:[10.1037/0033-295x.102.3.419](https://doi.org/10.1037/0033-295x.102.3.419). URL <http://dx.doi.org/10.1037/0033-295x.102.3.419>.
- Michael McCloskey and Neal J. Cohen. *Catastrophic Interference in Connectionist Networks: The Sequential Learning Problem*, pages 109–165. Elsevier, 1989. ISBN 9780125433242. doi:[10.1016/s0079-7421\(08\)60536-8](https://doi.org/10.1016/s0079-7421(08)60536-8). URL [http://dx.doi.org/10.1016/s0079-7421\(08\)60536-8](http://dx.doi.org/10.1016/s0079-7421(08)60536-8).
- Mark D. McDonnell and Lawrence M. Ward. The benefits of noise in neural systems: bridging theory and experiment. *Nature Reviews Neuroscience*, 12(7):415–425, 2011. ISSN 1471-0048. doi:[10.1038/nrn3061](https://doi.org/10.1038/nrn3061). URL <http://dx.doi.org/10.1038/nrn3061>.
- Denis Melanson, Mohammad Abu Khater, Maxwell Aifer, Kaelan Donatella, Max Hunter Gordon, Thomas Ahle, Gavin Crooks, Antonio J. Martinez, Faris Sbahi, and Patrick J. Coles. Thermodynamic computing system for AI applications. *Nature Communications*, 16(1), apr 22 2025. ISSN 2041-1723. doi:[10.1038/s41467-025-59011-x](https://doi.org/10.1038/s41467-025-59011-x). URL <http://dx.doi.org/10.1038/s41467-025-59011-x>.
- Leenoy Meshulam, Jeffrey L. Gauthier, Carlos D. Brody, David W. Tank, and William Bialek. Coarse Graining, Fixed Points, and Scaling in a Large Population of Neurons. *Physical Review Letters*, 123(17), 10 2019. ISSN 1079-7114. doi:[10.1103/physrevlett.123.178103](https://doi.org/10.1103/physrevlett.123.178103). URL <http://dx.doi.org/10.1103/physrevlett.123.178103>.
- Mohammed R. Milad and Gregory J. Quirk. Fear Extinction as a Model for Translational Neuroscience: Ten Years of Progress. *Annual Review of Psychology*, 63(1):129–151, jan 10 2012. ISSN 0066-4308. doi:[10.1146/annurev.psych.121208.131631](https://doi.org/10.1146/annurev.psych.121208.131631). URL <http://dx.doi.org/10.1146/annurev.psych.121208.131631>.
- Beren Millidge, Alexander Tschantz, and Christopher L. Buckley. Whence the Expected Free Energy? *Neural Computation*, 33(2):447–482, 2 2021. ISSN 1530-888X. doi:[10.1162/neco_a_01354](https://doi.org/10.1162/neco_a_01354). URL http://dx.doi.org/10.1162/neco_a_01354.
- Beren Millidge, Yuhang Song, Tommaso Salvatori, Thomas Lukasiewicz, and Rafal Bogacz. A Theoretical Framework for Inference and Learning in Predictive Coding Networks, 2022. URL <https://arxiv.org/abs/2207.12316>.
- MATHIEU MOLITOR. Remarks ON THE STATISTICAL ORIGIN OF THE GEOMETRICAL FORMULATION OF QUANTUM MECHANICS. *International Journal of Geometric Methods in Modern Physics*, 09(03):1220001, 5 2012. ISSN 1793-6977. doi:[10.1142/s0219887812200010](https://doi.org/10.1142/s0219887812200010). URL <http://dx.doi.org/10.1142/s0219887812200010>.
- Mathieu Molitor. Exponential families, Kähler geometry and quantum mechanics. *Journal of Geometry and Physics*, 70:54–80, 8 2013. ISSN 0393-0440. doi:[10.1016/j.geomphys.2013.03.015](https://doi.org/10.1016/j.geomphys.2013.03.015). URL <http://dx.doi.org/10.1016/j.geomphys.2013.03.015>.
- Alexander Mordvintsev, Ettore Randazzo, Eyvind Niklasson, and Michael Levin. Growing Neural Cellular Automata. *Distill*, 5(2), 2 2020. ISSN 2476-0757. doi:[10.23915/distill.00023](https://doi.org/10.23915/distill.00023). URL <http://dx.doi.org/10.23915/distill.00023>.
- Hazime Mori. Transport, Collective Motion, and Brownian Motion. *Progress of Theoretical Physics*, 33(3):423–455, 3 1965. ISSN 0033-068X. doi:[10.1143/ptp.33.423](https://doi.org/10.1143/ptp.33.423). URL <http://dx.doi.org/10.1143/ptp.33.423>.

- Bennet B. Murdock. The serial position effect of free recall. *Journal of Experimental Psychology*, 64(5):482–488, 11 1962. ISSN 0022-1015. doi:[10.1037/h0045106](https://doi.org/10.1037/h0045106). URL <http://dx.doi.org/10.1037/h0045106>.
- John D Murray, Alberto Bernacchia, David J Freedman, Ranulfo Romo, Jonathan D Wallis, Xinying Cai, Camillo Padoa-Schioppa, Tatiana Pasternak, Hyojung Seo, Daeyeol Lee, and Xiao-Jing Wang. A hierarchy of intrinsic timescales across primate cortex. *Nature Neuroscience*, 17(12):1661–1663, 11 2014. ISSN 1546-1726. doi:[10.1038/nn.3862](https://doi.org/10.1038/nn.3862). URL <http://dx.doi.org/10.1038/nn.3862>.
- J. P Nadal, G Toulouse, J. P Changeux, and S Dehaene. Networks of Formal Neurons and Memory Palimpsests. *Europhysics Letters (EPL)*, 2(4):343–343, aug 15 1986. ISSN 0295-5075. doi:[10.1209/0295-5075/2/4/013](https://doi.org/10.1209/0295-5075/2/4/013). URL <http://dx.doi.org/10.1209/0295-5075/2/4/013>.
- Erkki Oja. Simplified neuron model as a principal component analyzer. *Journal of Mathematical Biology*, 15(3):267–273, 11 1982. ISSN 1432-1416. doi:[10.1007/bf00275687](https://doi.org/10.1007/bf00275687). URL <http://dx.doi.org/10.1007/bf00275687>.
- Lars Onsager. Reciprocal Relations in Irreversible Processes. II. *Physical Review*, 38(12):2265–2279, 12 1931. ISSN 0031-899X. doi:[10.1103/physrev.38.2265](https://doi.org/10.1103/physrev.38.2265). URL <http://dx.doi.org/10.1103/physrev.38.2265>.
- Gergő Orbán, Pietro Berkes, József Fiser, and Máté Lengyel. Neural Variability and Sampling-Based Probabilistic Representations in the Visual Cortex. *Neuron*, 92(2):530–543, 10 2016. ISSN 0896-6273. doi:[10.1016/j.neuron.2016.09.038](https://doi.org/10.1016/j.neuron.2016.09.038). URL <http://dx.doi.org/10.1016/j.neuron.2016.09.038>.
- Pierre-Yves Oudeyer. What is intrinsic motivation? A typology of computational approaches. *Frontiers in Neurorobotics*, 1, 2007. ISSN 1662-5218. doi:[10.3389/neuro.12.006.2007](https://doi.org/10.3389/neuro.12.006.2007). URL <http://dx.doi.org/10.3389/neuro.12.006.2007>.
- Enzor Rafael Palacios, Adeel Razi, Thomas Parr, Michael Kirchhoff, and Karl Friston. On Markov blankets and hierarchical self-organisation. *Journal of Theoretical Biology*, 486:110089, 2 2020. ISSN 0022-5193. doi:[10.1016/j.jtbi.2019.110089](https://doi.org/10.1016/j.jtbi.2019.110089). URL <http://dx.doi.org/10.1016/j.jtbi.2019.110089>.
- G Parisi. A memory which forgets. *Journal of Physics A: Mathematical and General*, 19(10):L617–L620, 1986. ISSN 1361-6447. doi:[10.1088/0305-4470/19/10/011](https://doi.org/10.1088/0305-4470/19/10/011). URL <http://dx.doi.org/10.1088/0305-4470/19/10/011>.
- Thomas Parr, Giovanni Pezzulo, and Karl J. Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. The MIT Press, 3 2022. ISBN 9780262369978. doi:[10.7551/mitpress/12441.001.0001](https://doi.org/10.7551/mitpress/12441.001.0001). URL <http://dx.doi.org/10.7551/mitpress/12441.001.0001>.
- Cengiz Pehlevan and Dmitri B. Chklovskii. Neuroscience-Inspired Online Unsupervised Learning Algorithms: Artificial Neural Networks. *IEEE Signal Processing Magazine*, 36(6):88–96, 11 2019. ISSN 1558-0792. doi:[10.1109/msp.2019.2933846](https://doi.org/10.1109/msp.2019.2933846). URL <http://dx.doi.org/10.1109/msp.2019.2933846>.
- L. Personnaz, I. Guyon, and G. Dreyfus. Collective computational properties of neural networks: New learning mechanisms. *Physical Review A*, 34(5):4217–4228, 11 1986. ISSN 0556-2791. doi:[10.1103/physreva.34.4217](https://doi.org/10.1103/physreva.34.4217). URL <http://dx.doi.org/10.1103/physreva.34.4217>.
- Brad E. Pfeiffer and David J. Foster. Hippocampal place-cell sequences depict future paths to remembered goals. *Nature*, 497(7447):74–79, 4 2013. ISSN 1476-4687. doi:[10.1038/nature12112](https://doi.org/10.1038/nature12112). URL <http://dx.doi.org/10.1038/nature12112>.
- Giovanni Pistone and Carlo Sempì. An Infinite-Dimensional Geometric Structure on the Space of all the Probability Measures Equivalent to a Given One. *The Annals of Statistics*, 23(5), 10 1995. ISSN 0090-5364. doi:[10.1214/aos/1176324311](https://doi.org/10.1214/aos/1176324311). URL <http://dx.doi.org/10.1214/aos/1176324311>.
- Giulia L Poerio and Theodoros Karapanagiotidis. The default mode network and the complex dynamics of ongoing experience: an attractor-state perspective. *Current Opinion in Behavioral Sciences*, 65:101546, 10 2025. ISSN 2352-1546. doi:[10.1016/j.cobeha.2025.101546](https://doi.org/10.1016/j.cobeha.2025.101546). URL <http://dx.doi.org/10.1016/j.cobeha.2025.101546>.
- Emmanuel M. Pothos and Jerome R. Busemeyer. Can quantum probability provide a new direction for cognitive modeling? *Behavioral and Brain Sciences*, 36(3):255–274, 5 2013. ISSN 1469-1825. doi:[10.1017/s0140525x12001525](https://doi.org/10.1017/s0140525x12001525). URL <http://dx.doi.org/10.1017/s0140525x12001525>.
- A. R. Powers, C. Mathys, and P. R. Corlett. Pavlovian conditioning-induced hallucinations result from overweighting of perceptual priors. *Science*, 357(6351):596–600, 8 2017. ISSN 1095-9203. doi:[10.1126/science.aan3458](https://doi.org/10.1126/science.aan3458). URL <http://dx.doi.org/10.1126/science.aan3458>.
- M. Prezioso, F. Merrih-Bayat, B. D. Hoskins, G. C. Adam, K. K. Likharev, and D. B. Strukov. Training and operation of an integrated neuromorphic network based on metal-oxide memristors. *Nature*, 521(7550):61–64, 5 2015. ISSN 1476-4687. doi:[10.1038/nature14441](https://doi.org/10.1038/nature14441). URL <http://dx.doi.org/10.1038/nature14441>.

- Misha Rabinovich, Ramon Huerta, and Gilles Laurent. Transient Dynamics for Neural Processing. *Science*, 321(5885):48–50, 2008. ISSN 1095-9203. doi:[10.1126/science.1155564](https://doi.org/10.1126/science.1155564). URL <http://dx.doi.org/10.1126/science.1155564>.
- Marcus E. Raichle, Ann Mary MacLeod, Abraham Z. Snyder, William J. Powers, Debra A. Gusnard, and Gordon L. Shulman. A default mode of brain function. *Proceedings of the National Academy of Sciences*, 98(2):676–682, 1 2001. ISSN 1091-6490. doi:[10.1073/pnas.98.2.676](https://doi.org/10.1073/pnas.98.2.676). URL <http://dx.doi.org/10.1073/pnas.98.2.676>.
- Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler, Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Victor Greiff, David Kreil, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield Networks is All You Need, 2020. URL <https://arxiv.org/abs/2008.02217>.
- Maxwell James Ramstead, Mahault Albarracin, Alex Kiefer, Brennan Klein, Chris Fields, Karl Friston, and Adam Safron. The inner screen model of consciousness: applying the free energy principle directly to the study of conscious experience. 4 2023. doi:[10.31234/osf.io/6afs3](https://doi.org/10.31234/osf.io/6afs3). URL <http://dx.doi.org/10.31234/osf.io/6afs3>.
- Rajesh P. N. Rao and Dana H. Ballard. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, 1 1999. ISSN 1546-1726. doi:[10.1038/4580](https://doi.org/10.1038/4580). URL <http://dx.doi.org/10.1038/4580>.
- A. David Redish, Steve Jensen, Adam Johnson, and Zeb Kurth-Nelson. Reconciling reinforcement learning models with behavioral extinction and renewal: Implications for addiction, relapse, and problem gambling. *Psychological Review*, 114(3):784–805, 2007. ISSN 0033-295X. doi:[10.1037/0033-295x.114.3.784](https://doi.org/10.1037/0033-295x.114.3.784). URL <http://dx.doi.org/10.1037/0033-295x.114.3.784>.
- Robert A. Rescorla. Spontaneous Recovery. *Learning & Memory*, 11(5):501–509, 2004. ISSN 1549-5485. doi:[10.1101/lm.77504](https://doi.org/10.1101/lm.77504). URL <http://dx.doi.org/10.1101/lm.77504>.
- Robert A. Rescorla and C. Donald Heth. Reinstatement of fear to an extinguished conditioned stimulus. *Journal of Experimental Psychology: Animal Behavior Processes*, 1(1):88–96, 1975. ISSN 0097-7403. doi:[10.1037/0097-7403.1.1.88](https://doi.org/10.1037/0097-7403.1.1.88). URL <http://dx.doi.org/10.1037/0097-7403.1.1.88>.
- Luc Rey-Bellet and Konstantinos Spiliopoulos. Irreversible Langevin samplers and variance reduction: a large deviations approach. *Nonlinearity*, 28(7):2081–2103, 5 2015. ISSN 1361-6544. doi:[10.1088/0951-7715/28/7/2081](https://doi.org/10.1088/0951-7715/28/7/2081). URL <http://dx.doi.org/10.1088/0951-7715/28/7/2081>.
- Blake A. Richards, Timothy P. Lillicrap, Philippe Beaudoin, Yoshua Bengio, Rafal Bogacz, Amelia Christensen, Claudia Clopath, Rui Ponte Costa, Archy de Berker, Surya Ganguli, Colleen J. Gillon, Daniyar Hafner, Adam Kepecs, Nikolaus Kriegeskorte, Peter Latham, Grace W. Lindsay, Kenneth D. Miller, Richard Naud, Christopher C. Pack, Panayiota Poirazi, Pieter Roelfsema, João Sacramento, Andrew Saxe, Benjamin Scellier, Anna C. Schapiro, Walter Senn, Greg Wayne, Daniel Yamins, Friedemann Zenke, Joel Zylberberg, Denis Therien, and Konrad P. Kording. A deep learning framework for neuroscience. *Nature Neuroscience*, 22(11):1761–1770, 10 2019. ISSN 1546-1726. doi:[10.1038/s41593-019-0520-2](https://doi.org/10.1038/s41593-019-0520-2). URL <http://dx.doi.org/10.1038/s41593-019-0520-2>.
- Fernando E. Rosas, Pedro A. M. Mediano, Henrik J. Jensen, Anil K. Seth, Adam B. Barrett, Robin L. Carhart-Harris, and Daniel Bor. Reconciling emergences: An information-theoretic approach to identify causal emergence in multivariate data. *PLOS Computational Biology*, 16(12):e1008289, 12 2020. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1008289](https://doi.org/10.1371/journal.pcbi.1008289). URL <http://dx.doi.org/10.1371/journal.pcbi.1008289>.
- Fernando E. Rosas, Bernhard C. Geiger, Andrea I. Luppi, Anil K. Seth, Daniel Polani, Michael Gastpar, and Pedro A. M. Mediano. Software in the natural world: A computational approach to hierarchical emergence, 2024. URL <https://arxiv.org/abs/2402.09090>.
- Fernando E. Rosas, Andrea I. Luppi, Pedro A. M. Mediano, Morten L. Kringelbach, Luiz Pessoa, and Federico Turkheimer. Top-down and bottom-up neuroscience: overcoming the clash of research cultures. *Nature Reviews Neuroscience*, 26(9):513–515, 2025. ISSN 1471-0048. doi:[10.1038/s41583-025-00946-x](https://doi.org/10.1038/s41583-025-00946-x). URL <http://dx.doi.org/10.1038/s41583-025-00946-x>.
- Claudio Runfola, Spase Petkoski, Hiba Sheheitli, Christophe Bernard, Anthony R. McIntosh, and Viktor Jirsa. A mechanism for the emergence of low-dimensional structures in brain dynamics. *npj Systems Biology and Applications*, 11(1), 4 2025. ISSN 2056-7189. doi:[10.1038/s41540-025-00499-w](https://doi.org/10.1038/s41540-025-00499-w). URL <http://dx.doi.org/10.1038/s41540-025-00499-w>.
- Dalton A R Sakthivadivel. Towards a Geometry and Analysis for Bayesian Mechanics, 2022. URL <https://arxiv.org/abs/2204.11900>.

- Paula Sanz Leon, Stuart A. Knock, M. Marmaduke Woodman, Lia Domide, Jochen Mersmann, Anthony R. McIntosh, and Viktor Jirsa. The Virtual Brain: a simulator of primate brain network dynamics. *Frontiers in Neuroinformatics*, 7, 2013. ISSN 1662-5196. doi:[10.3389/fninf.2013.00010](https://doi.org/10.3389/fninf.2013.00010). URL <http://dx.doi.org/10.3389/fninf.2013.00010>.
- Benjamin Scellier and Yoshua Bengio. Equilibrium Propagation: Bridging the Gap between Energy-Based Models and Backpropagation. *Frontiers in Computational Neuroscience*, 11, 5 2017. ISSN 1662-5188. doi:[10.3389/fncom.2017.00024](https://doi.org/10.3389/fncom.2017.00024). URL <http://dx.doi.org/10.3389/fncom.2017.00024>.
- Jürgen Schmidhuber. Formal Theory of Creativity, Fun, and Intrinsic Motivation (1990–2010). *IEEE Transactions on Autonomous Mental Development*, 2(3):230–247, 2010. ISSN 1943-0612. doi:[10.1109/tamd.2010.2056368](https://doi.org/10.1109/tamd.2010.2056368). URL <http://dx.doi.org/10.1109/tamd.2010.2056368>.
- J. Schnakenberg. Network theory of microscopic and macroscopic behavior of master equation systems. *Reviews of Modern Physics*, 48(4):571–585, 10 1976. ISSN 0034-6861. doi:[10.1103/revmodphys.48.571](https://doi.org/10.1103/revmodphys.48.571). URL <http://dx.doi.org/10.1103/revmodphys.48.571>.
- Elad Schneidman, Michael J. Berry, Ronen Segev, and William Bialek. Weak pairwise correlations imply strongly correlated network states in a neural population. *Nature*, 440(7087):1007–1012, 4 2006. ISSN 1476-4687. doi:[10.1038/nature04701](https://doi.org/10.1038/nature04701). URL <http://dx.doi.org/10.1038/nature04701>.
- Catherine D. Schuman, Shruti R. Kulkarni, Maryam Parsa, J. Parker Mitchell, Prasanna Date, and Bill Kay. Opportunities for neuromorphic computing algorithms and applications. *Nature Computational Science*, 2(1):10–19, 1 2022. ISSN 2662-8457. doi:[10.1038/s43588-021-00184-y](https://doi.org/10.1038/s43588-021-00184-y). URL <http://dx.doi.org/10.1038/s43588-021-00184-y>.
- Thomas Seidenbecher, T. Rao Laxmi, Oliver Stork, and Hans-Christian Pape. Amygdalar and Hippocampal Theta Rhythm Synchronization During Fear Memory Retrieval. *Science*, 301(5634):846–850, 8 2003. ISSN 1095-9203. doi:[10.1126/science.1085818](https://doi.org/10.1126/science.1085818). URL <http://dx.doi.org/10.1126/science.1085818>.
- Walter Senn, Dominik Dold, Akos F Kungl, Benjamin Ellenberger, Jakob Jordan, Yoshua Bengio, João Sacramento, and Mihai A Petrovici. A neuronal least-action principle for real-time learning in cortical circuits. *eLife*, 12, 12 2024. ISSN 2050-084X. doi:[10.7554/elife.89674](https://doi.org/10.7554/elife.89674). URL <http://dx.doi.org/10.7554/eLife.89674>.
- Ludovica Serricchio, Dario Bocchi, Claudio Chilin, Raffaele Marino, Matteo Negri, Chiara Cammarota, and Federico Ricci-Tersenghi. Daydreaming Hopfield Networks and their surprising effectiveness on correlated data. *Neural Networks*, 186:107216, 2025. ISSN 0893-6080. doi:[10.1016/j.neunet.2025.107216](https://doi.org/10.1016/j.neunet.2025.107216). URL <http://dx.doi.org/10.1016/j.neunet.2025.107216>.
- Anil K. Seth, Adam B. Barrett, and Lionel Barnett. Granger Causality Analysis in Neuroscience and Neuroimaging. *The Journal of Neuroscience*, 35(8):3293–3297, 2 2015. ISSN 1529-2401. doi:[10.1523/jneurosci.4399-14.2015](https://doi.org/10.1523/jneurosci.4399-14.2015). URL <http://dx.doi.org/10.1523/jneurosci.4399-14.2015>.
- Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. Continual Learning with Deep Generative Replay, 2017. URL <https://arxiv.org/abs/1705.08690>.
- Herbert A. Simon. The architecture of complexity. *Proceedings of the American Philosophical Society*, 106(6):467–482, 1962.
- Matthew F. Singh, Todd S. Braver, Michael W. Cole, and ShiNung Ching. Estimation and validation of individualized dynamic brain models with resting state fMRI. *NeuroImage*, 221:117046, 11 2020. ISSN 1053-8119. doi:[10.1016/j.neuroimage.2020.117046](https://doi.org/10.1016/j.neuroimage.2020.117046). URL <http://dx.doi.org/10.1016/j.neuroimage.2020.117046>.
- N. A Sinitsyn and I Nemenman. The Berry phase and the pump flux in stochastic chemical kinetics. *Europhysics Letters (EPL)*, 77(5):58001, 2 2007. ISSN 1286-4854. doi:[10.1209/0295-5075/77/58001](https://doi.org/10.1209/0295-5075/77/58001). URL <http://dx.doi.org/10.1209/0295-5075/77/58001>.
- Stephen M. Smith, Peter T. Fox, Karla L. Miller, David C. Glahn, P. Mickle Fox, Clare E. Mackay, Nicola Filippini, Kate E. Watkins, Roberto Toro, Angela R. Laird, and Christian F. Beckmann. Correspondence of the brain’s functional architecture during activation and rest. *Proceedings of the National Academy of Sciences*, 106(31):13040–13045, 8 2009. ISSN 1091-6490. doi:[10.1073/pnas.0905267106](https://doi.org/10.1073/pnas.0905267106). URL <http://dx.doi.org/10.1073/pnas.0905267106>.
- Stephen M. Smith, Karla L. Miller, Gholamreza Salimi-Khorshidi, Matthew Webster, Christian F. Beckmann, Thomas E. Nichols, Joseph D. Ramsey, and Mark W. Woolrich. Network modelling methods for FMRI. *NeuroImage*, 54(2):875–891, 1 2011. ISSN 1053-8119. doi:[10.1016/j.neuroimage.2010.08.063](https://doi.org/10.1016/j.neuroimage.2010.08.063). URL <http://dx.doi.org/10.1016/j.neuroimage.2010.08.063>.

- Paul Smolensky. *Information processing in dynamical systems: foundations of harmony theory*, volume 1, pages 194–281. MIT Press, 1986.
- H. Sompolinsky and I. Kanter. Temporal Association in Asymmetric Neural Networks. *Physical Review Letters*, 57(22):2861–2864, 12 1986. ISSN 0031-9007. doi:[10.1103/physrevlett.57.2861](https://doi.org/10.1103/physrevlett.57.2861). URL <http://dx.doi.org/10.1103/physrevlett.57.2861>.
- Yuhang Song, Beren Millidge, Tommaso Salvatori, Thomas Lukasiewicz, Zhenghua Xu, and Rafal Bogacz. Inferring neural activity before plasticity as a foundation for learning beyond backpropagation. *Nature Neuroscience*, 27(2):348–358, 1 2024. ISSN 1546-1726. doi:[10.1038/s41593-023-01514-1](https://doi.org/10.1038/s41593-023-01514-1). URL <http://dx.doi.org/10.1038/s41593-023-01514-1>.
- Tamas Spisak and Karl Friston. Self-orthogonalizing attractor neural networks emerging from the free energy principle. *Neurocomputing*, 682:133472, 2026. ISSN 0925-2312. doi:[10.1016/j.neucom.2026.133472](https://doi.org/10.1016/j.neucom.2026.133472). URL <http://dx.doi.org/10.1016/j.neucom.2026.133472>.
- Tamas Spisak, Balint Kincses, Frederik Schlitt, Matthias Zunhammer, Tobias Schmidt-Wilcke, Zsigmond T. Kincses, and Ulrike Bingel. Pain-free resting-state functional brain connectivity predicts individual pain sensitivity. *Nature Communications*, 11(1), 1 2020. ISSN 2041-1723. doi:[10.1038/s41467-019-13785-z](https://doi.org/10.1038/s41467-019-13785-z). URL <http://dx.doi.org/10.1038/s41467-019-13785-z>.
- Tamas Spisak, Ulrike Bingel, and Tor D. Wager. Multivariate BWAS can be replicable with moderate sample sizes. *Nature*, 615(7951):E4–E7, 3 2023. ISSN 1476-4687. doi:[10.1038/s41586-023-05745-x](https://doi.org/10.1038/s41586-023-05745-x). URL <http://dx.doi.org/10.1038/s41586-023-05745-x>.
- Kimberly L Stachenfeld, Matthew M Botvinick, and Samuel J Gershman. The hippocampus as a predictive map. *Nature Neuroscience*, 20(11):1643–1653, 10 2017. ISSN 1546-1726. doi:[10.1038/nn.4650](https://doi.org/10.1038/nn.4650). URL <http://dx.doi.org/10.1038/nn.4650>.
- Menachem Stern and Arvind Murugan. Learning Without Neurons in Physical Systems. *Annual Review of Condensed Matter Physics*, 14(1):417–441, 3 2023. ISSN 1947-5462. doi:[10.1146/annurev-conmatphys-040821-113439](https://doi.org/10.1146/annurev-conmatphys-040821-113439). URL <http://dx.doi.org/10.1146/annurev-conmatphys-040821-113439>.
- Menachem Stern, Daniel Hexner, Jason W. Rocks, and Andrea J. Liu. Supervised Learning in Physical Networks: From Machine Learning to Learning Machines. *Physical Review X*, 11(2), 5 2021. ISSN 2160-3308. doi:[10.1103/physrevx.11.021045](https://doi.org/10.1103/physrevx.11.021045). URL <http://dx.doi.org/10.1103/physrevx.11.021045>.
- Menachem Stern, Marcelo Guzman, Felipe Martins, Andrea J. Liu, and Vijay Balasubramanian. Physical Networks Become What They Learn. *Physical Review Letters*, 134(14), 4 2025. ISSN 1079-7114. doi:[10.1103/physrevlett.134.147402](https://doi.org/10.1103/physrevlett.134.147402). URL <http://dx.doi.org/10.1103/physrevlett.134.147402>.
- Susanne Still, David A. Sivak, Anthony J. Bell, and Gavin E. Crooks. Thermodynamics of Prediction. *Physical Review Letters*, 109(12), 2012. ISSN 1079-7114. doi:[10.1103/physrevlett.109.120604](https://doi.org/10.1103/physrevlett.109.120604). URL <http://dx.doi.org/10.1103/physrevlett.109.120604>.
- Mark G. Stokes. ‘Activity-silent’ working memory in prefrontal cortex: a dynamic coding framework. *Trends in Cognitive Sciences*, 19(7):394–405, 2015. ISSN 1364-6613. doi:[10.1016/j.tics.2015.05.004](https://doi.org/10.1016/j.tics.2015.05.004). URL <http://dx.doi.org/10.1016/j.tics.2015.05.004>.
- Dmitri B. Strukov, Gregory S. Snider, Duncan R. Stewart, and R. Stanley Williams. The missing memristor found. *Nature*, 453(7191):80–83, 5 2008. ISSN 1476-4687. doi:[10.1038/nature06932](https://doi.org/10.1038/nature06932). URL <http://dx.doi.org/10.1038/nature06932>.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.
- Abby Tabor, Michael A. Thacker, G. Lorimer Moseley, and Konrad P. Körding. Pain: A Statistical Account. *PLOS Computational Biology*, 13(1):e1005142, 1 2017. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1005142](https://doi.org/10.1371/journal.pcbi.1005142). URL <http://dx.doi.org/10.1371/journal.pcbi.1005142>.
- Emmanuelle Tognoli and J. A. Scott Kelso. The Metastable Brain. *Neuron*, 81(1):35–48, 1 2014. ISSN 0896-6273. doi:[10.1016/j.neuron.2013.12.022](https://doi.org/10.1016/j.neuron.2013.12.022). URL <http://dx.doi.org/10.1016/j.neuron.2013.12.022>.
- Mark K. Transtrum, Benjamin B. Machta, Kevin S. Brown, Bryan C. Daniels, Christopher R. Myers, and James P. Sethna. Perspective: Sloppiness and emergent theories in physics, biology, and beyond. *The Journal of Chemical Physics*, 143(1), 2015. ISSN 1089-7690. doi:[10.1063/1.4923066](https://doi.org/10.1063/1.4923066). URL <http://dx.doi.org/10.1063/1.4923066>.
- Ichiro Tsuda. Toward an interpretation of dynamic neural activity in terms of chaotic dynamical systems. *Behavioral and Brain Sciences*, 24(5):793–810, 10 2001. ISSN 1469-1825. doi:[10.1017/s0140525x01000097](https://doi.org/10.1017/s0140525x01000097). URL <http://dx.doi.org/10.1017/s0140525x01000097>.

- Sebastian Urchs, Jonathan Armoza, Clara Moreau, Yassine Benhajali, Jolène St-Aubin, Pierre Orban, and Pierre Bellec. Mist: A multi-resolution parcellation of functional brain networks. *MNI Open Research*, 1:3, 3 2019. ISSN 2515-5059. doi:[10.12688/mniopenres.12767.2](https://doi.org/10.12688/mniopenres.12767.2). URL <http://dx.doi.org/10.12688/mniopenres.12767.2>.
- Gido M. van de Ven, Hava T. Siegelmann, and Andreas S. Tolias. Brain-inspired replay for continual learning with artificial neural networks. *Nature Communications*, 11(1), 8 2020. ISSN 2041-1723. doi:[10.1038/s41467-020-17866-2](https://doi.org/10.1038/s41467-020-17866-2). URL <http://dx.doi.org/10.1038/s41467-020-17866-2>.
- Vitaly Vanchurin, Yuri I. Wolf, Mikhail I. Katsnelson, and Eugene V. Koonin. Toward a theory of evolution as multilevel learning. *Proceedings of the National Academy of Sciences*, 119(6), 2 2022. ISSN 1091-6490. doi:[10.1073/pnas.2120037119](https://doi.org/10.1073/pnas.2120037119). URL <http://dx.doi.org/10.1073/pnas.2120037119>.
- Diego Vidaurre, Stephen M. Smith, and Mark W. Woolrich. Brain network dynamics are hierarchically organized in time. *Proceedings of the National Academy of Sciences*, 114(48):12827–12832, 10 2017. ISSN 1091-6490. doi:[10.1073/pnas.1705120114](https://doi.org/10.1073/pnas.1705120114). URL <http://dx.doi.org/10.1073/pnas.1705120114>.
- Tor D. Wager and Lauren Y. Atlas. The neuroscience of placebo effects: connecting context, learning and health. *Nature Reviews Neuroscience*, 16(7):403–418, 2015. ISSN 1471-0048. doi:[10.1038/nrn3976](https://doi.org/10.1038/nrn3976). URL <http://dx.doi.org/10.1038/nrn3976>.
- Jin Wang, Li Xu, and Erkang Wang. Potential landscape and flux framework of nonequilibrium networks: Robustness, dissipation, and coherence of biochemical oscillations. *Proceedings of the National Academy of Sciences*, 105(34):12271–12276, 8 2008. ISSN 1091-6490. doi:[10.1073/pnas.0800579105](https://doi.org/10.1073/pnas.0800579105). URL <http://dx.doi.org/10.1073/pnas.0800579105>.
- Xiao-Jing Wang. Synaptic reverberation underlying mnemonic persistent activity. *Trends in Neurosciences*, 24(8):455–463, 8 2001. ISSN 0166-2236. doi:[10.1016/s0166-2236\(00\)01868-3](https://doi.org/10.1016/s0166-2236(00)01868-3). URL [http://dx.doi.org/10.1016/s0166-2236\(00\)01868-3](http://dx.doi.org/10.1016/s0166-2236(00)01868-3).
- Takamitsu Watanabe, Satoshi Hirose, Hiroyuki Wada, Yoshio Imai, Toru Machida, Ichiro Shirouzu, Seiki Konishi, Yasushi Miyashita, and Naoki Masuda. A pairwise maximum entropy model accurately describes resting-state human brain networks. *Nature Communications*, 4(1), 1 2013. ISSN 2041-1723. doi:[10.1038/ncomms2388](https://doi.org/10.1038/ncomms2388). URL <http://dx.doi.org/10.1038/ncomms2388>.
- Richard Watson and Michael Levin. The collective intelligence of evolution and development. *Collective Intelligence*, 2(2):263391372311683, 4 2023. ISSN 2633-9137. doi:[10.1177/26339137231168355](https://doi.org/10.1177/26339137231168355). URL <http://dx.doi.org/10.1177/26339137231168355>.
- Richard A. Watson, Rob Mills, and C. L. Buckley. Global Adaptation in Networks of Selfish Components: Emergent Associative Memory at the System Scale. *Artificial Life*, 17(3):147–166, 2011. ISSN 1530-9185. doi:[10.1162/artl_a_00029](https://doi.org/10.1162/artl_a_00029). URL http://dx.doi.org/10.1162/artl_a_00029.
- James C. R. Whittington and Rafal Bogacz. An Approximation of the Error Backpropagation Algorithm in a Predictive Coding Network with Local Hebbian Synaptic Plasticity. *Neural Computation*, 29(5):1229–1262, may 1 2017. ISSN 0899-7667. doi:[10.1162/neco_a_00949](https://doi.org/10.1162/neco_a_00949). URL http://dx.doi.org/10.1162/neco_a_00949.
- Christopher J. Whyte and Ryan Smith. The predictive global neuronal workspace: A formal active inference model of visual consciousness. *Progress in Neurobiology*, 199:101918, 4 2021. ISSN 0301-0082. doi:[10.1016/j.pneurobio.2020.101918](https://doi.org/10.1016/j.pneurobio.2020.101918). URL <http://dx.doi.org/10.1016/j.pneurobio.2020.101918>.
- Matthew A. Wilson and Bruce L. McNaughton. Reactivation of Hippocampal Ensemble Memories During Sleep. *Science*, 265(5172):676–679, jul 29 1994. ISSN 0036-8075. doi:[10.1126/science.8036517](https://doi.org/10.1126/science.8036517). URL <http://dx.doi.org/10.1126/science.8036517>.
- Stephen Wolfram. Universality and complexity in cellular automata. *Physica D: Nonlinear Phenomena*, 10(1-2):1–35, 1 1984. ISSN 0167-2789. doi:[10.1016/0167-2789\(84\)90245-8](https://doi.org/10.1016/0167-2789(84)90245-8). URL [http://dx.doi.org/10.1016/0167-2789\(84\)90245-8](http://dx.doi.org/10.1016/0167-2789(84)90245-8).
- James Woodward. *Making Things Happen: A Theory of Causal Explanation*. Oxford University Press New York, 1 2004. ISBN 9780199835089. doi:[10.1093/0195155270.001.0001](https://doi.org/10.1093/0195155270.001.0001). URL <http://dx.doi.org/10.1093/0195155270.001.0001>.
- Yuichi Yamashita and Jun Tani. Emergence of Functional Hierarchy in a Multiple Timescale Neural Network Model: A Humanoid Robot Experiment. *PLoS Computational Biology*, 4(11):e1000220, 11 2008. ISSN 1553-7358. doi:[10.1371/journal.pcbi.1000220](https://doi.org/10.1371/journal.pcbi.1000220). URL <http://dx.doi.org/10.1371/journal.pcbi.1000220>.

- Maria Yampolskaya and Pankaj Mehta. Hopfield Networks as Models of Emergent Function in Biology. *Annual Review of Biophysics*, 55(1):323–342, 5 2026. ISSN 1936-1238. doi:[10.1146/annurev-biophys-101425-023356](https://doi.org/10.1146/annurev-biophys-101425-023356). URL <http://dx.doi.org/10.1146/annurev-biophys-101425-023356>.
- Han Yan, Lei Zhao, Liang Hu, Xidi Wang, Erkang Wang, and Jin Wang. Nonequilibrium landscape theory of neural networks. *Proceedings of the National Academy of Sciences*, 110(45), 10 2013. ISSN 1091-6490. doi:[10.1073/pnas.1310692110](https://doi.org/10.1073/pnas.1310692110). URL <http://dx.doi.org/10.1073/pnas.1310692110>.
- Anthony Zador, Sean Escola, Blake Richards, Bence Ölveczky, Yoshua Bengio, Kwabena Boahen, Matthew Botvinick, Dmitri Chklovskii, Anne Churchland, Claudia Clopath, James DiCarlo, Surya Ganguli, Jeff Hawkins, Konrad Körding, Alexei Koulakov, Yann LeCun, Timothy Lillicrap, Adam Marblestone, Bruno Olshausen, Alexandre Pouget, Cristina Savin, Terrence Sejnowski, Eero Simoncelli, Sara Solla, David Sussillo, Andreas S. Tolias, and Doris Tsao. Catalyzing next-generation Artificial Intelligence through NeuroAI. *Nature Communications*, 14(1), 3 2023. ISSN 2041-1723. doi:[10.1038/s41467-023-37180-x](https://doi.org/10.1038/s41467-023-37180-x). URL <http://dx.doi.org/10.1038/s41467-023-37180-x>.
- Guanxiong Zeng, Yang Chen, Bo Cui, and Shan Yu. Continual learning of context-dependent processing in neural networks. *Nature Machine Intelligence*, 1(8):364–372, 8 2019. ISSN 2522-5839. doi:[10.1038/s42256-019-0080-x](https://doi.org/10.1038/s42256-019-0080-x). URL <http://dx.doi.org/10.1038/s42256-019-0080-x>.
- Robert Zwanzig. Memory Effects in Irreversible Thermodynamics. *Physical Review*, 124(4):983–992, 11 1961. ISSN 0031-899X. doi:[10.1103/physrev.124.983](https://doi.org/10.1103/physrev.124.983). URL <http://dx.doi.org/10.1103/physrev.124.983>.